

Inteligencia Artificial no Google Cloud Platform

AI Platform, KubeFlow e TFX

Pythian



Recifense
MSc em Inteligencia Artificial

Data Scientist | ML Engineer @ Pythian
GDE in ML
GDG Cloud Ottawa-Montreal

 carlos-timoteo-data-scientist



Pythian



AI Platform Overview

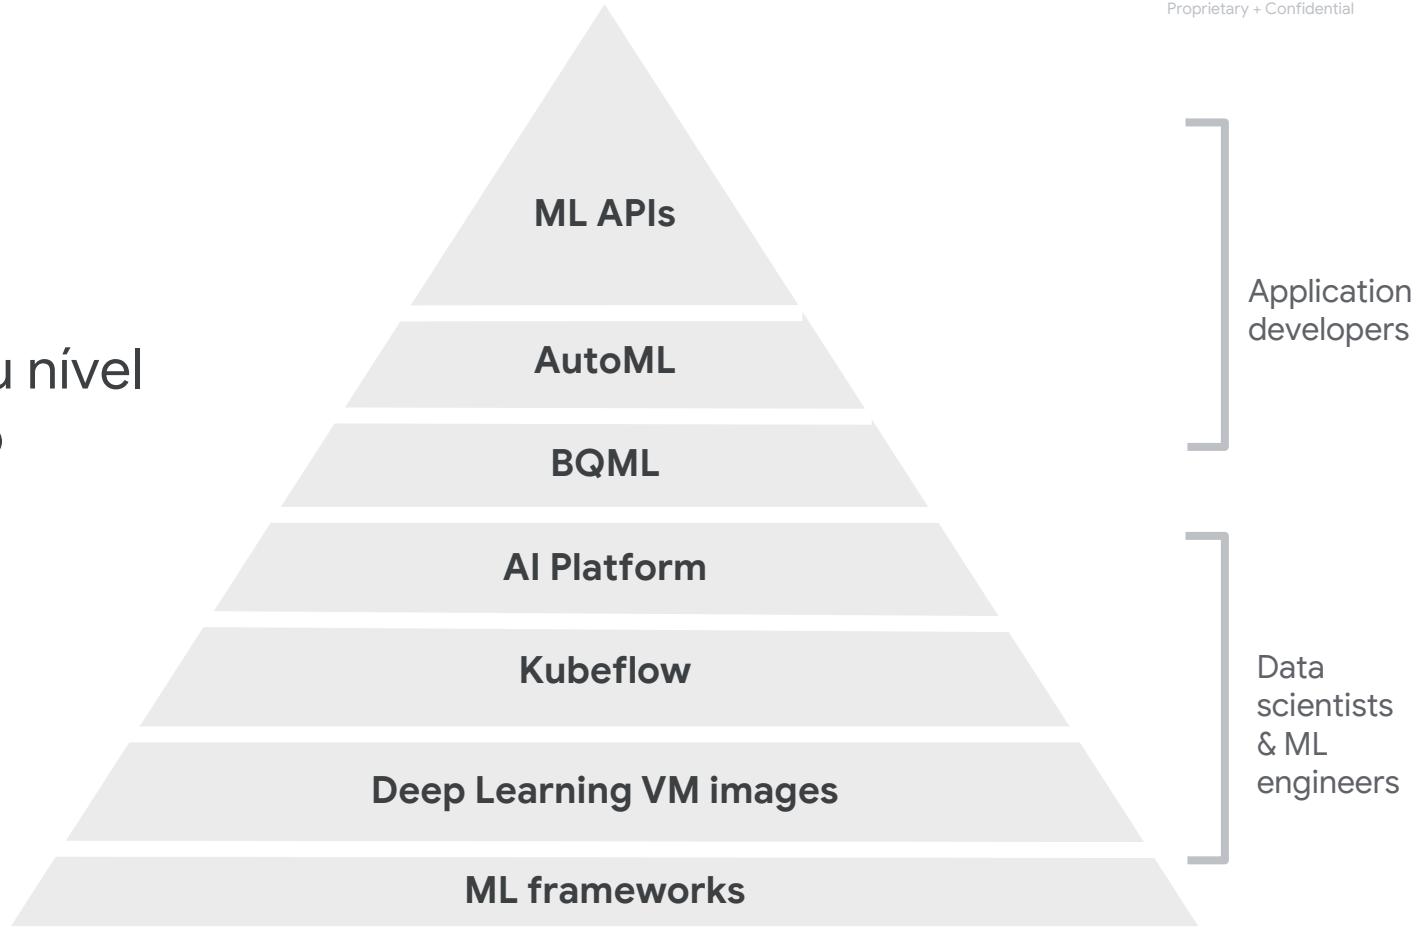


Google Cloud

Como conectar usuários para produzir resultados **10x** mais impactantes?

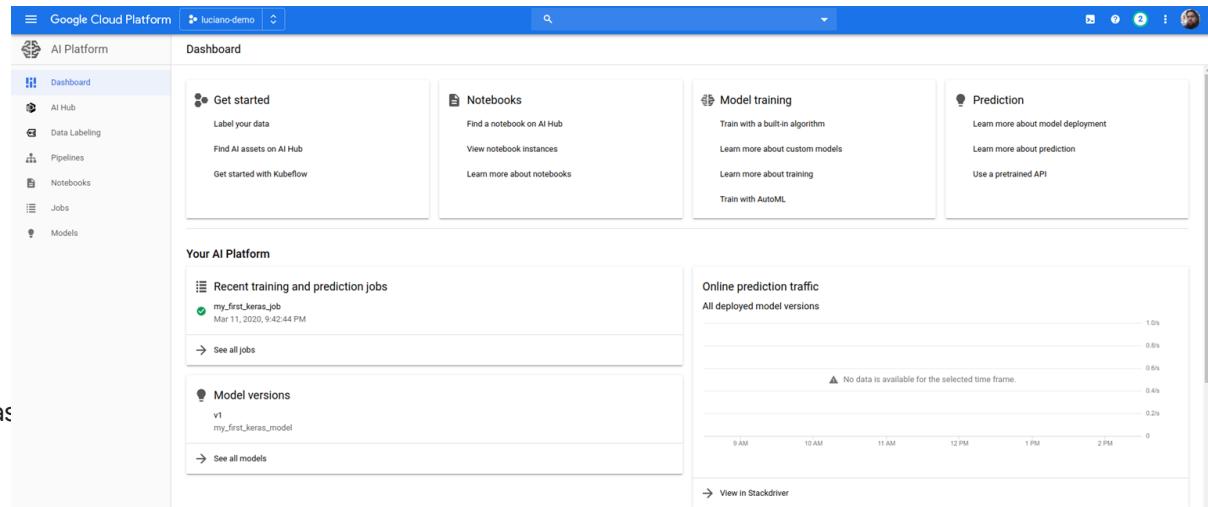


Escolha o seu nível
de abstração

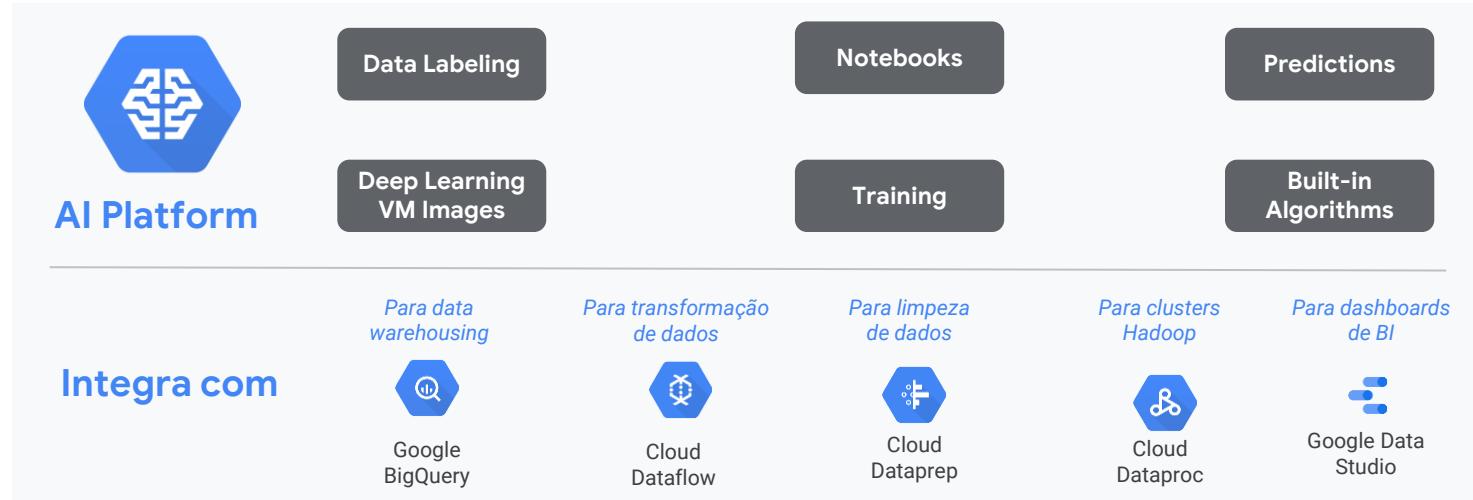


Introdução ao Cloud AI Platform

- Ambiente de desenvolvimento fim-a-fim para IA dentro da console GCP
- Construído com Kubeflow, oferece uma cadeia de ferramentas integrada - de engenharia de dados à implementação de modelos. Sem “lock-in”
- Permite implementar suas soluções on-premises ou na Google Cloud sem mudanças significativas de código
- Acesso à tecnologias Google AI de ponta como Tensorflow, Tensorflow Extended (TFX), TPUs quando implementando sua solução em produção



O que está incluso no Cloud AI Platform



AI Hub



Kubeflow
(On premises)



Deep learning VM image (DLVM)

Machine Learning mais rápido e fácil no GCE

- **Prototipação Acelerada**

Protótipo seu projeto rapidamente utilizando VMs pré-configuradas para Deep Learning

- **Suporte a CPU, GPU e TPU**

Utilize aceleradores para seus modelos, como GPU ou TPU, com poucos cliques

- **Performance otimizada para Google Cloud**

Bibliotecas e configurações otimizadas para a melhor performance para a sua infraestrutura - assim você não se preocupa com isso

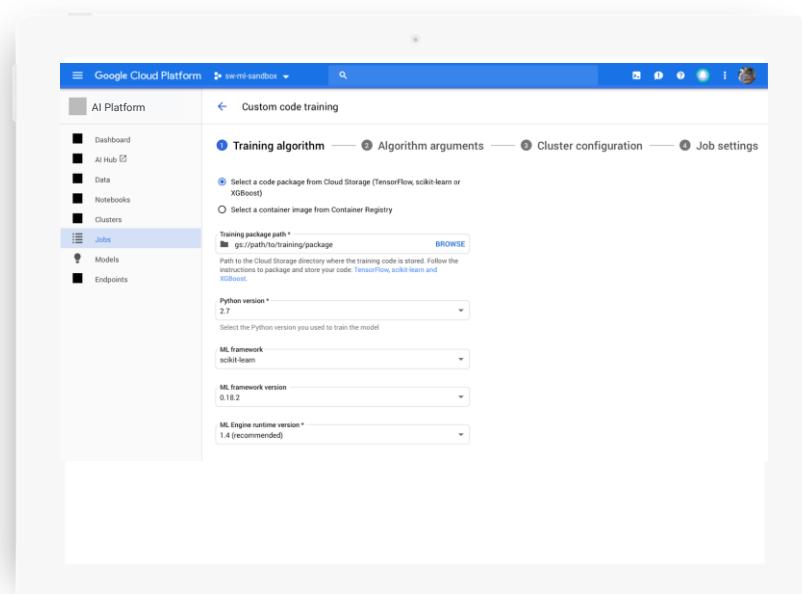
- **Flexibilidade**

Escolha entre diferentes ML frameworks como TensorFlow, PyTorch e scikit-learn ou instale o framework que você preferir sob a imagem base



Gerenciamento serverless de modelos

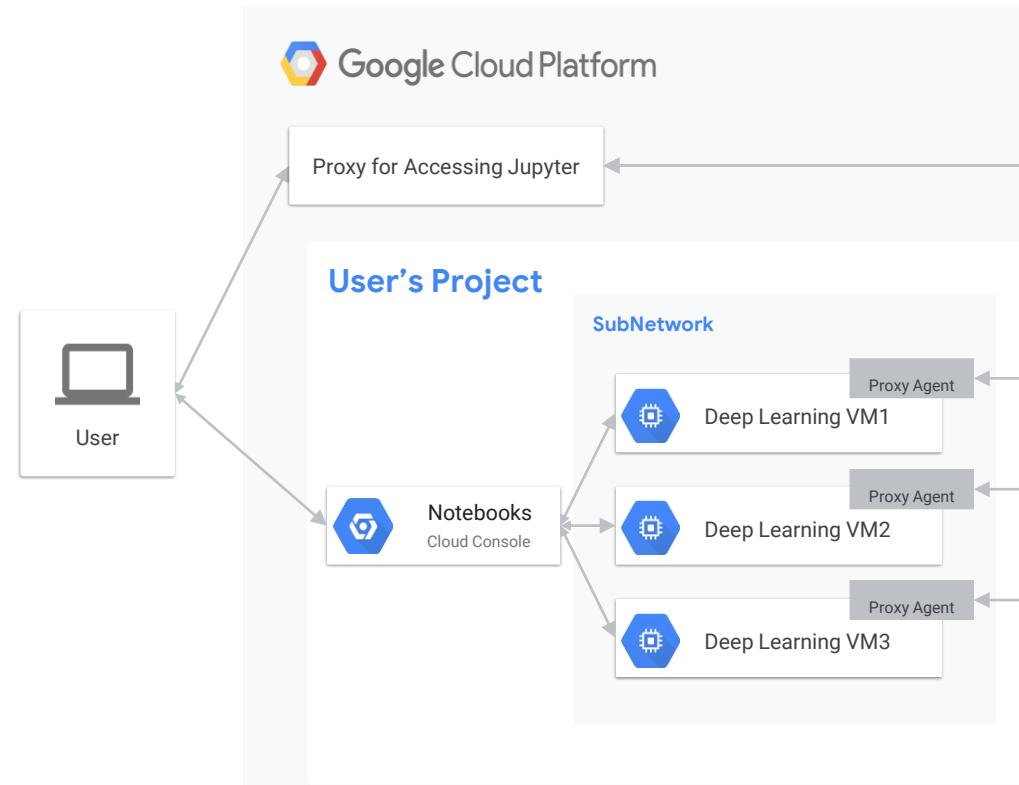
- Treine modelos abstraindo a infraestrutura
- Suporta os frameworks de machine learning mais populares. Suporta, também, containers Docker customizados
- Realiza treinamentos distribuídos e utiliza GPUs e TPUs para finalizar jobs mais rápido
- Melhora modelos através da otimização de hiperparâmetros automatizada



Notebooks gerenciados e ambientes pré-configurados

Bibliotecas GCP pré-instaladas

Ambientes pré-configurados para:



Crie uma instância de Notebook

Notebook instances BETA

+ NEW INSTANCE C REFRESH ▶ START ■ STOP ⚡ RESET 🗑 DELETE

Filter table

Instance name	Region	Framework	Machine type	GPUs	Labels
<input type="checkbox"/> cancer-demo-instance	us-west1-b	TensorFlow:1.13	4 vCPUs, 15 GB RAM	NVIDIA Tesla V100 x 1	No labels

Tensorflow ➔ Standard
us-west1-b, 4 vCPUs, 15 GB RAM, 100 GB Disk

Pytorch ➔ With GPU
us-west1-b, 4 vCPUs, 15 GB RAM, 1 NVIDIA Tesla K80, 100 GB Disk

More options



Notebook instances BETA

+ NEW INSTANCE C REFRESH ▶ START ■ STOP ⚡ RESET 🗑 DELETE

Filter table

Instance name	Region	Framework	Machine type	GPUs	Labels
<input checked="" type="checkbox"/> cancer-demo-instance	us-west1-b	TensorFlow:1.13	4 vCPUs, 15 GB RAM	NVIDIA Tesla V100 x 1	No labels

[OPEN JUPYTERLAB](#)

Importe seus dados do Cloud Storage

```
[ ]: !pip3 install tensorflow_hub  
  
[1]: import sys  
# Workaround: pip installs to a wrong location on these notebooks  
sys.path.insert(0,'/home/jupyter/.local/lib/python3.5/site-packages')  
import os  
import json  
import random  
import base64  
import tensorflow as tf  
import tensorflow_hub as hub  
import tensorflow.keras as ks  
import IPython.display  
from tensorflow.python.lib.io import file_io  
  
storage_bucket = 'gs://cancer-demo-data/ht-kg-histopathologic-cancer-detection/'  
train_files = [storage_bucket+f for f in ['train_0.tfrecords', 'train_1.tfrecords', 'train_2.tfrecords']]  
eval_files = [storage_bucket+f for f in ['train_3.tfrecords']]  
xval_files = [storage_bucket+f for f in ['train_4.tfrecords']]  
  
row_count_file = storage_bucket + 'record_counts.json'  
examples_file = storage_bucket + 'examples.zip'
```

Treine seu modelo onde quiser

On-Premises

```
from fairing.config.config_kubeflow import KubeflowConfig
from fairing.jobs.train_job import TrainJob

job = TrainJob(train,
                base_docker_image="gcr.io/caip-dexter-dev/jaas:py3.5.3",
                destination_dir='gs://cancer_detection_demo/',
                runtime_config=KubeflowConfig())

job.submit()
job.get_logs()
```

Google Cloud

```
from fairing.config.config_cmle import CMLEConfig
from fairing.jobs.train_job import TrainJob

job = TrainJob(train,
                base_docker_image="gcr.io/caip-dexter-dev/jaas:py3.5.3",
                destination_dir='gs://cancer_detection_demo/',
                runtime_config=CMLEConfig())

job.submit()
```

Simples modificações para treinar seu modelo on-premises com Kubeflow ou na Google Cloud using AI Platform

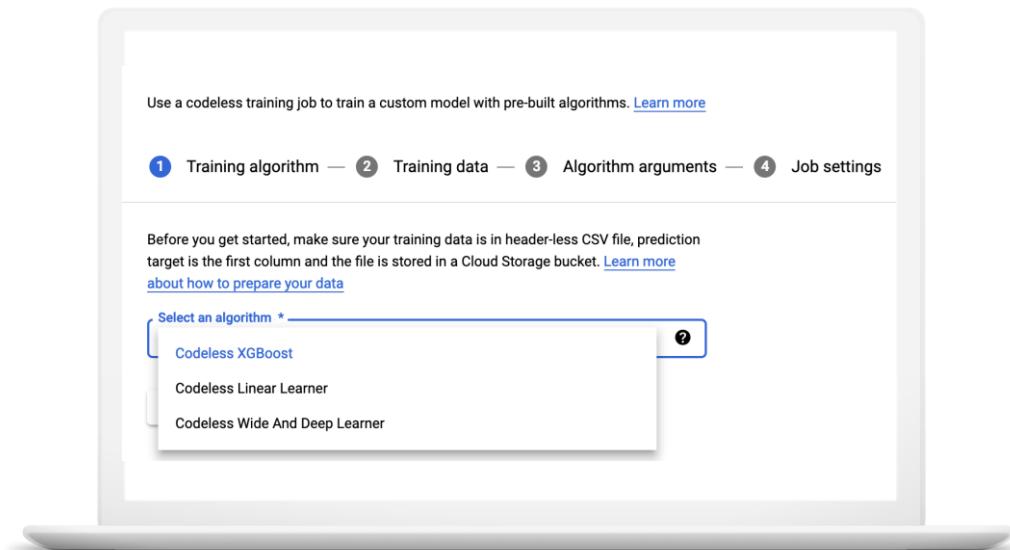
Implemente seu modelo com facilidade

```
%%bash -s "$SAVED_MODEL_DIR"
sm_dir=$1
model_name=kg_cancer_detection
version_name=versions

# Create the model and base version if it doesn't exist
gcloud ml-engine models describe $model_name > /dev/null
if [ $? -ne 0 ]; then gcloud ml-engine models create $model_name; fi
gcloud ml-engine versions describe $version_name --model $model_name 2> /dev/null
if [ $? -ne 0 ]; then
    gcloud ml-engine versions create $version_name \
        --model $model_name \
        --origin=$sm_dir \
fi
```

Modelos Built-in no Cloud AI Platform

- Otimização de hiperparâmetros embutida
- Suporte à algoritmos populares
- Busque soluções no AI Hub
- Facilidade para incluir novos algoritmos como containers



Três caminhos de AI na Google Cloud

seus dados + seu modelo



Cloud TPUs



Compute Engine



Cloud Dataproc



Kubernetes Engine



Cloud ML Engine



BigQuery ML

seus dados +
nossos modelos

AutoML



nossos dados + nossos modelos



Cloud
Translation API



Cloud
Vision API



Cloud
Speech API



Cloud
Video
Intelligence API



Data Loss
Prevention API



Cloud Speech
Synthesis API



Cloud Natural
Language API



Dialogflow

Customização



Construa seus modelos

Treine os nossos modelos

Use nossas APIs inteligentes

Cloud Cloud





BigQuery ML Overview



Google Cloud



BigQuery ML

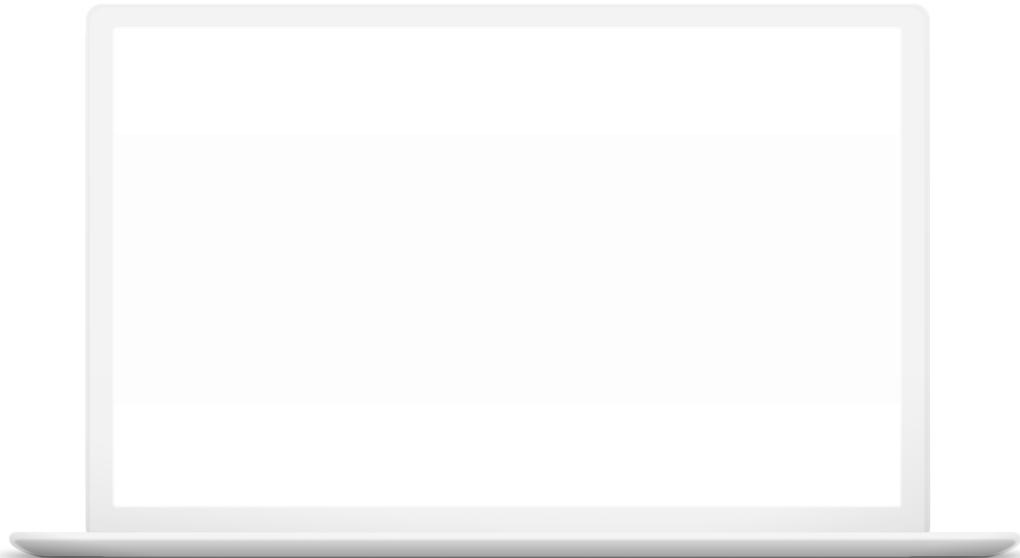
BigQuery ML
empowers both
data analysts and
data scientists

Execute ML initiatives without
moving data from BigQuery

Iterate on models in SQL in BigQuery to
increase development speed

Automate model selection, and hypertuning

BigQuery ML



- 1
- 2
- 3

Execute ML initiatives without moving data from BigQuery

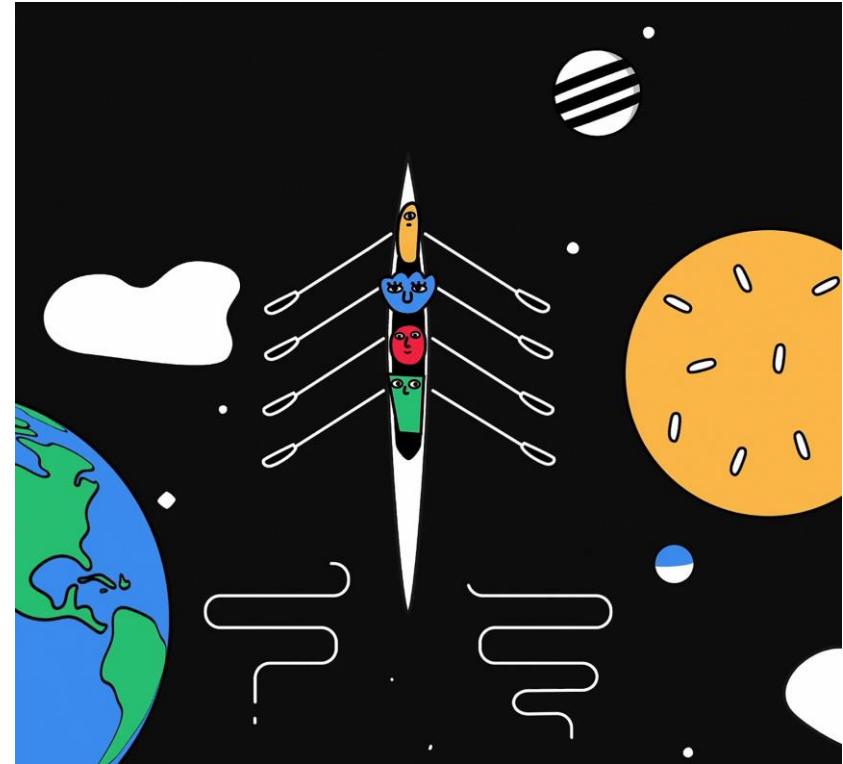
Iterate on models in SQL in BigQuery to increase development speed

Automate common ML tasks and hyperparameter tuning

Behind the scenes

Through two lines of SQL

- Leverage BigQuery's processing power to build a model
- Auto-tuned learning rate
- Auto-split of data into training and test
- Null imputation
- Standardization of numeric features
- One-hot encoding of strings
- Class imbalance handling



Hardware Acceleration

Cloud TPU v3

Accelerator families comparison



TPU v1

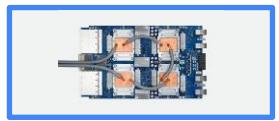
int8 only - inference only required model quantization



TPU v2

bfloat16 - inference and training no model changes necessary

180 TFLOPS (bfloat16)
/ board



TPU v3

bfloat16 - inference and training no model changes necessary

420 TFLOPS
(bfloat16)
/ board



NVIDIA P100

3584 CUDA cores - float16
but float16 x float16 => float32 not available

18 TFLOPS (float16)
/ chip



NVIDIA V100

5120 CUDA cores + 640 TensorCores - float16
model tweaking necessary for float16

112 TFLOPS (float16)
/ chip

Reference models for Cloud TPUs

Image Recognition
Object Detection



Image Recognition

AmoebaNet-D
ResNet-50/101/152/200
Inception v2/v3/v4
DenseNet

Object Detection

RetinaNet

Low-Resource Models

MobileNet
SqueezeNet

Machine Translation
& Language Modeling



Models

Machine translation
Language modeling
Sentiment analysis
Question-answering
(all Transformer-based)

Speech
Recognition



Models

ASR Transformer
(LibriSpeech)

Image
Generation



Models

Image Transformer
DCGAN

Training hardware options on AI Platform

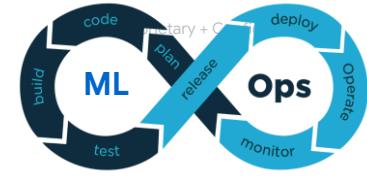
	Training Time	Cost
GPU - P100	5h50	\$15
GPU - V100	4h30	\$18
Cluster 5 GPUs - P100	1h15	\$15
Cluster 5 GPUs - V100	1h00	\$18
VM with 4 GPUs - V100	1h50	\$27
Cloud TPU v2	1h00	\$6.6

```
# config.yaml  
  
trainingInput:  
  scaleTier: BASIC_TPU
```

CLOUD TPU

8 CORES

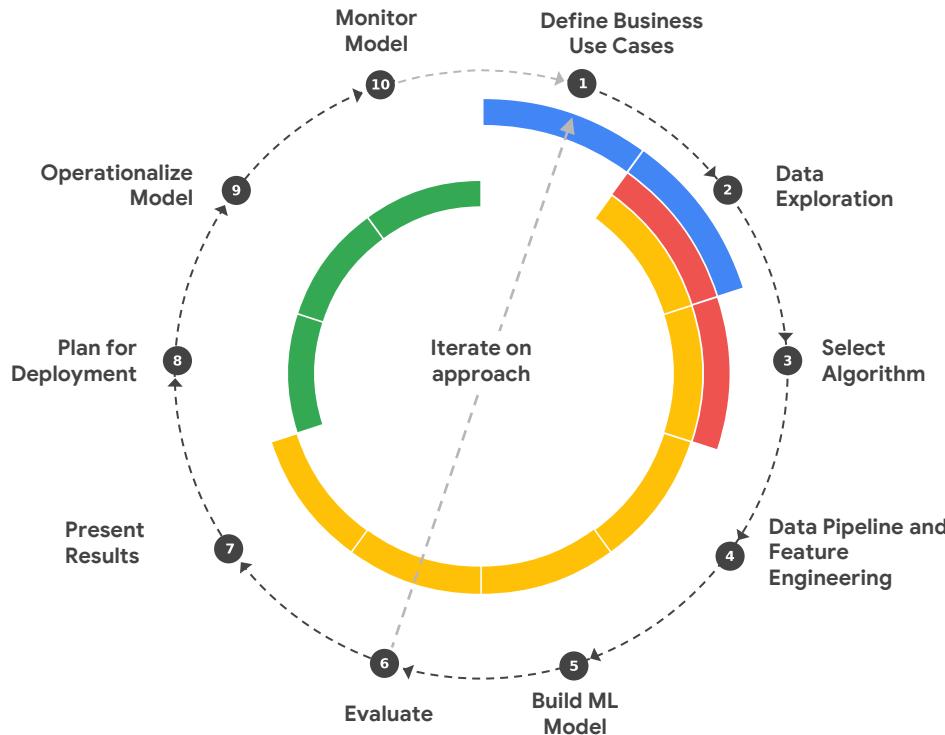




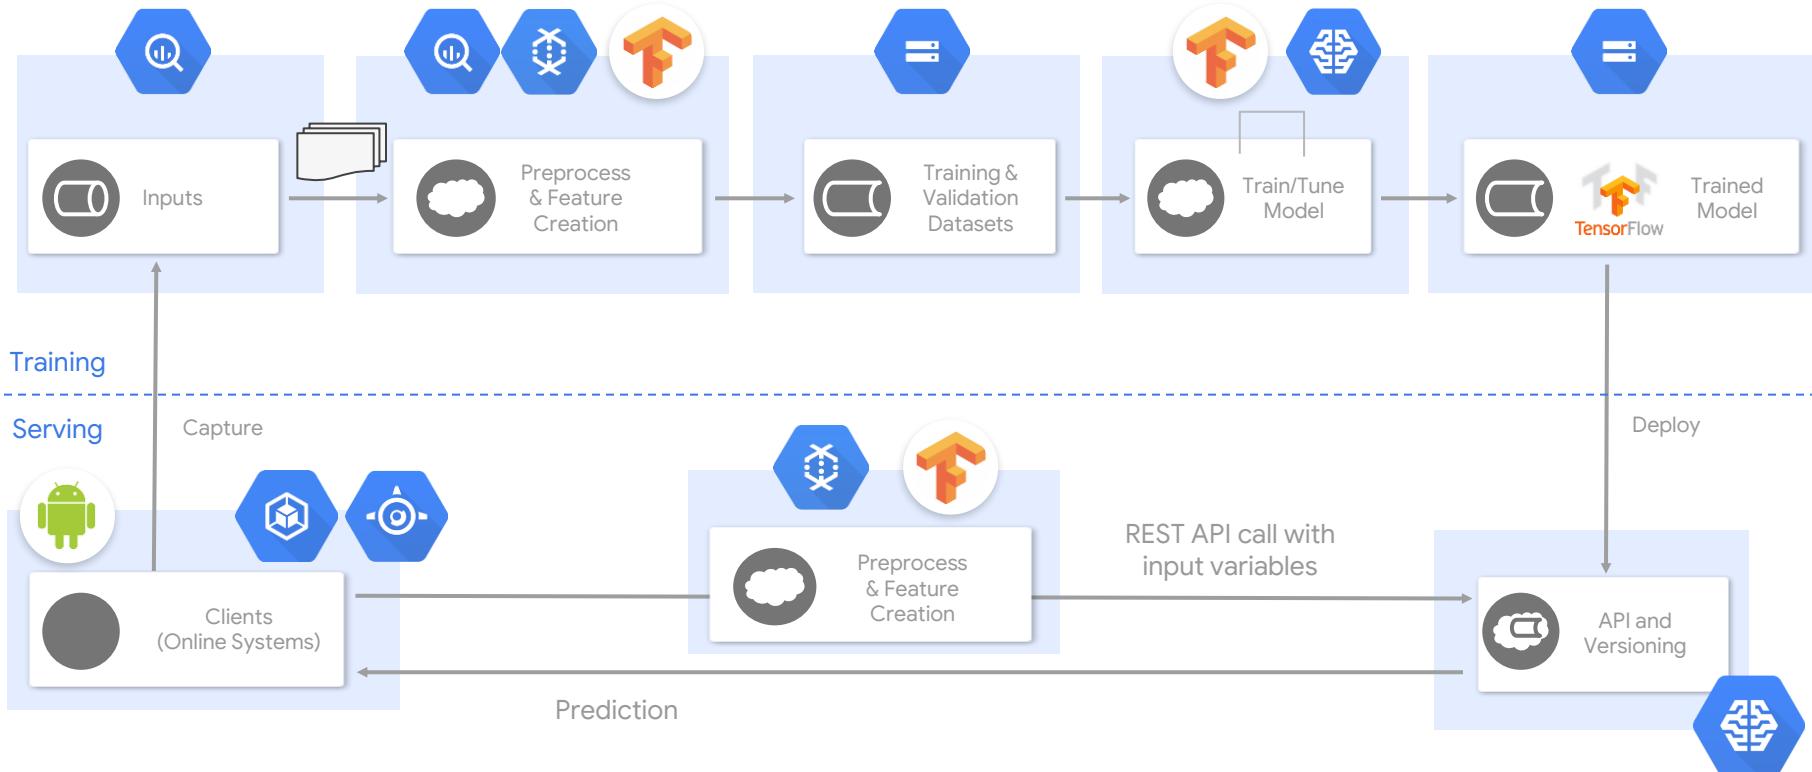
MLOps (ML Operationalization)

Typical machine learning lifecycle

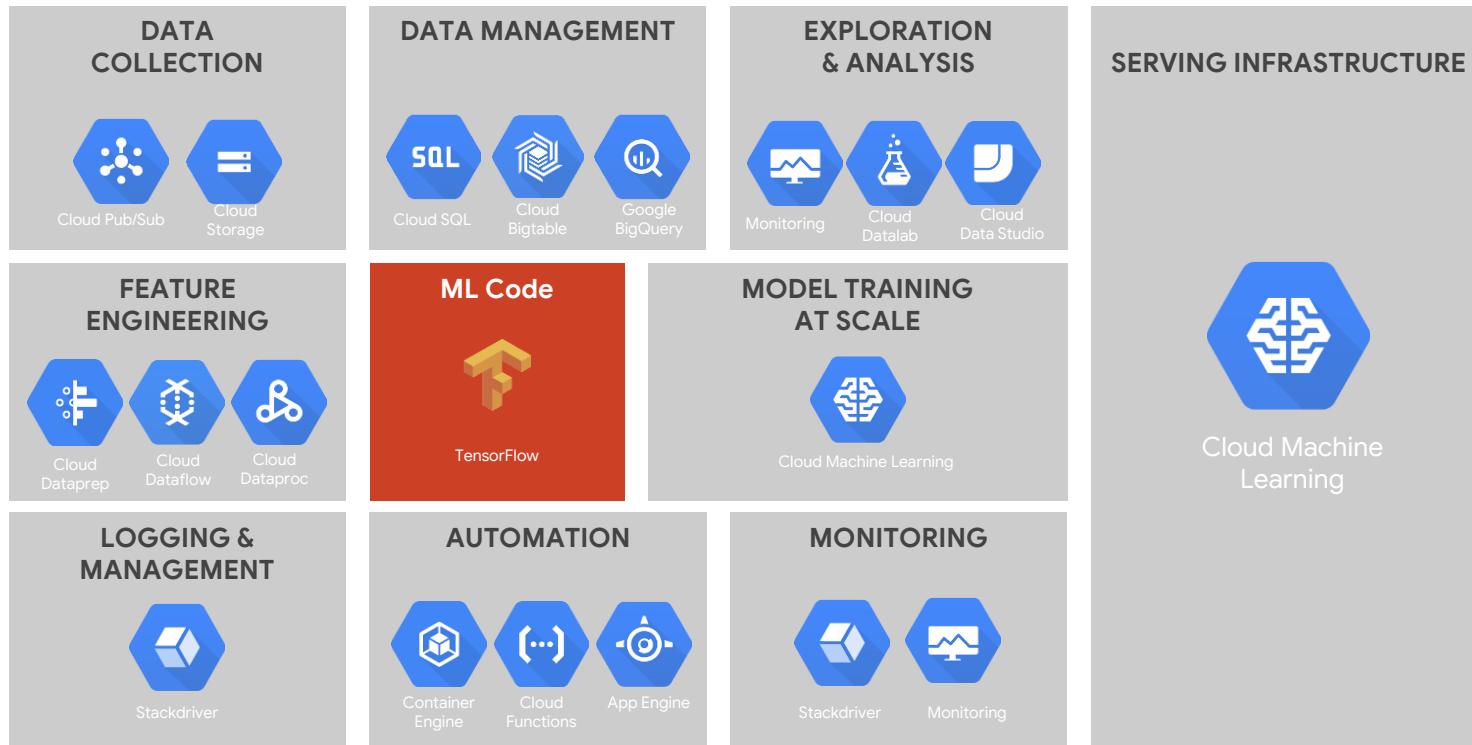
Step-by-step solution of ML problem



Machine learning pipeline @ GCP



Operational ML - end-to-end ML solution on GCP



We need cloud-native ML!



Composability

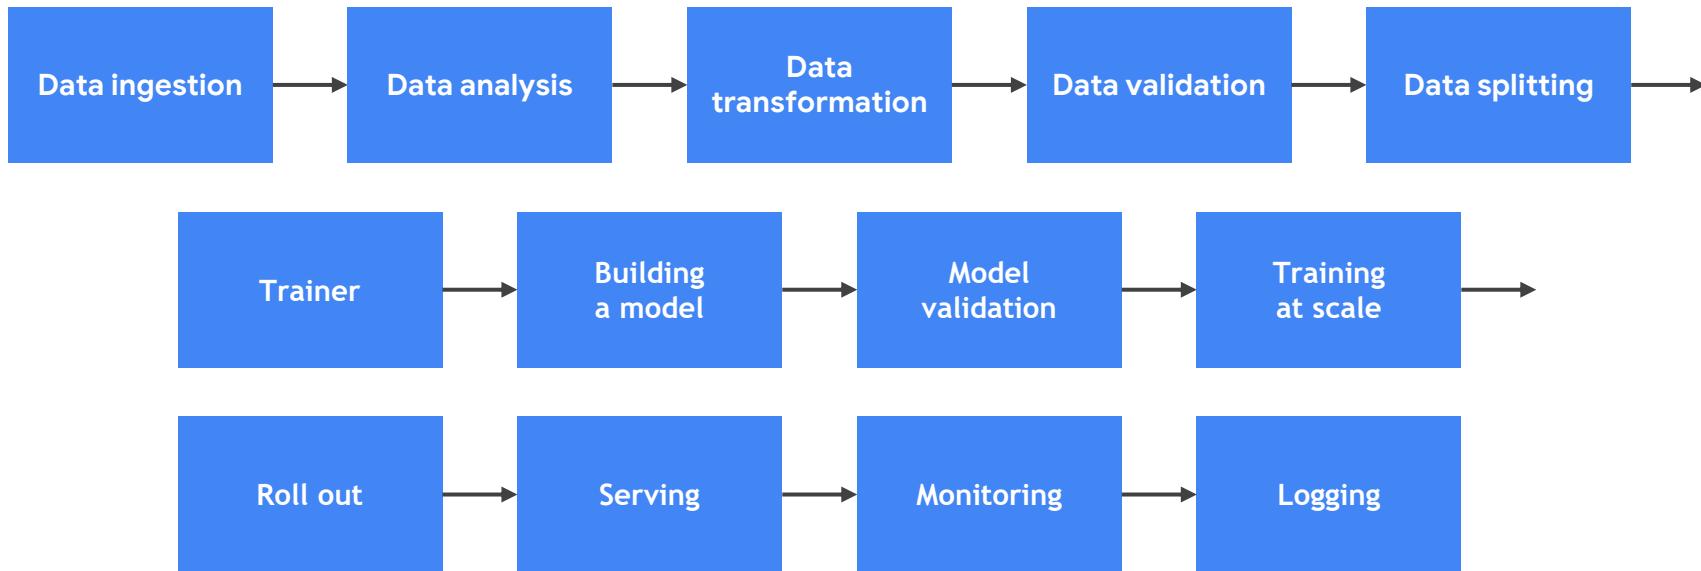


Portability

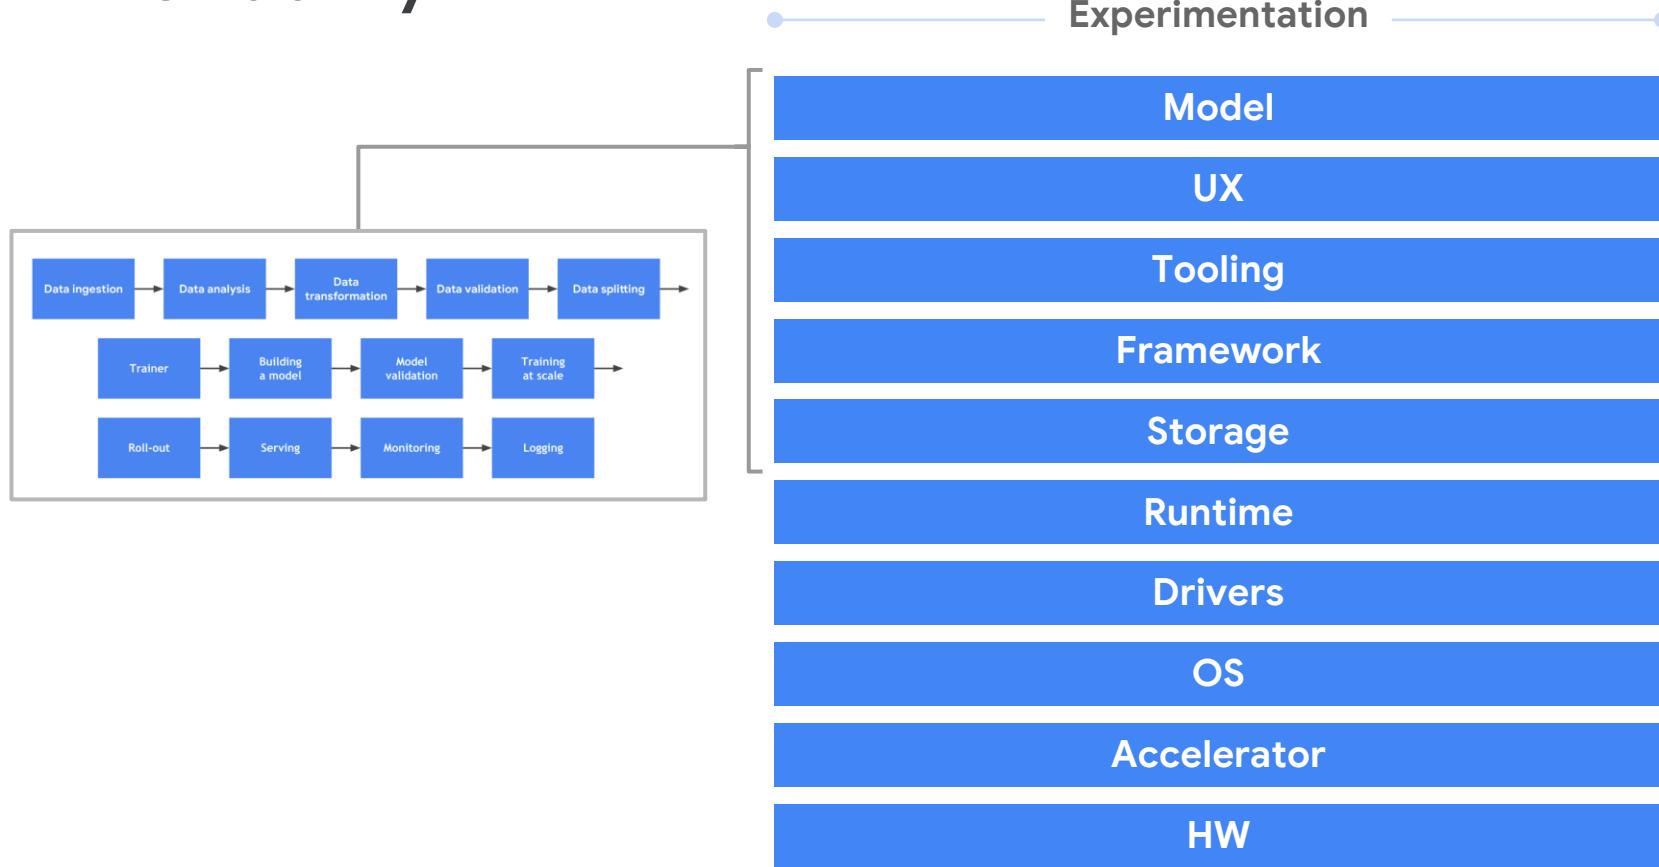


Scalability

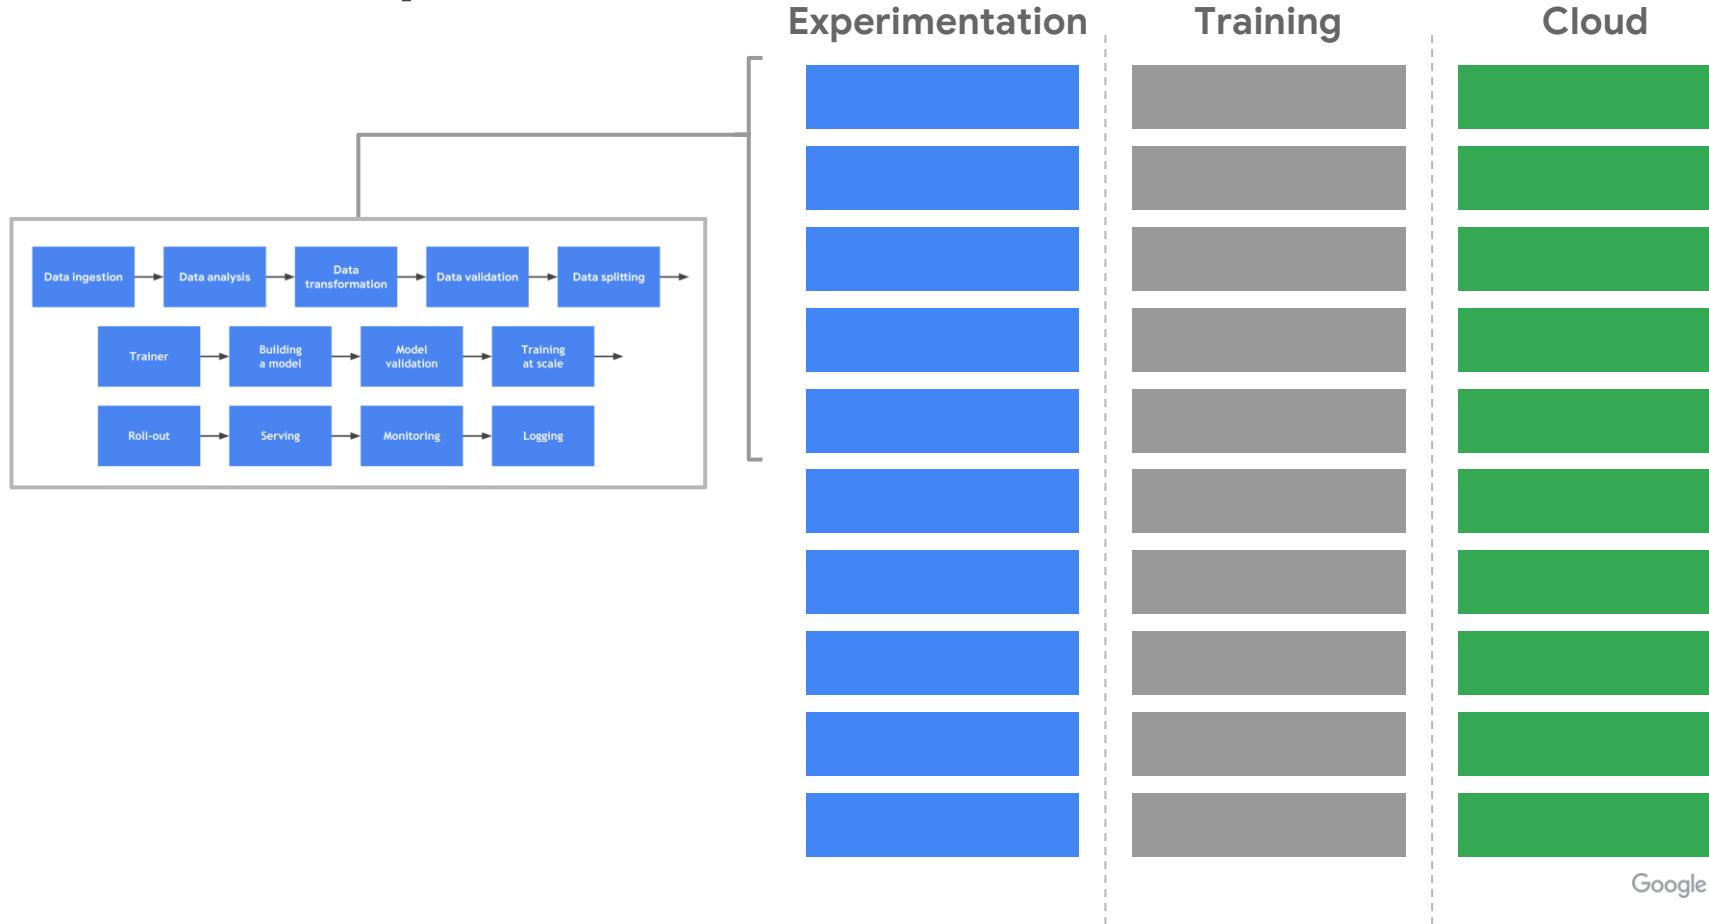
Composability



Portability



Portability



You know what's
really good at
composability,
portability and
scalability?



Containers & Kubernetes

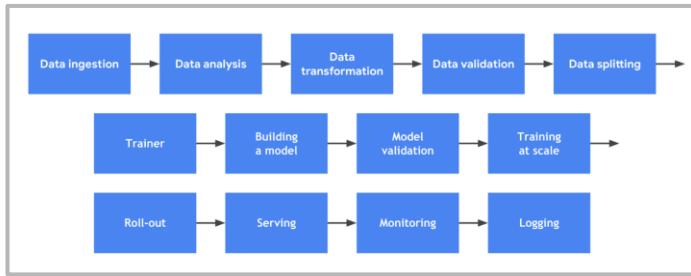
Oh, you want to use ML on Kubernetes?

First, become an expert in:

- Containers
- Packaging
- Kubernetes service endpoints
- Persistent volumes
- Scaling
- Immutable deployments
- GPUs, drivers, and the GPL
- Cloud APIs
- DevOps
- ...



Portability



Experimentation



Training



Cloud



Portability



Experimentation



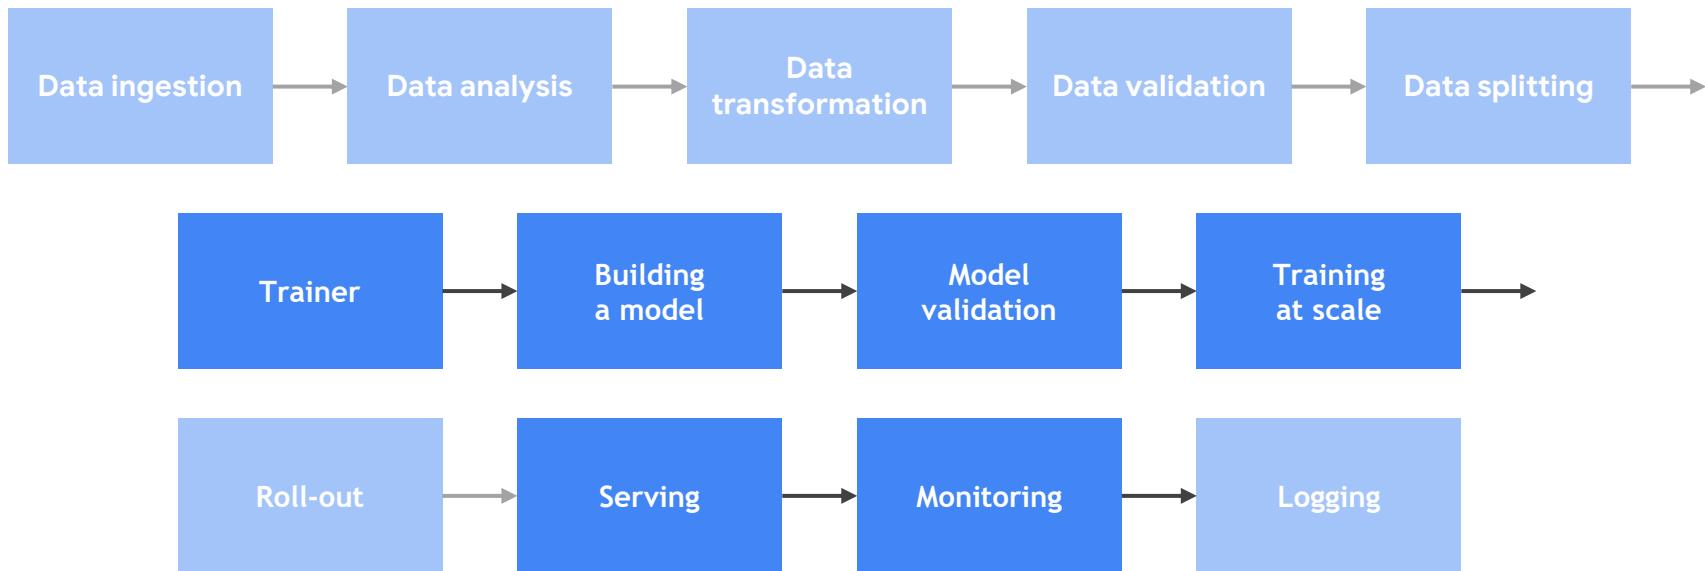
Training



Cloud



What's in the box?



Kubeflow is open!



Open
community



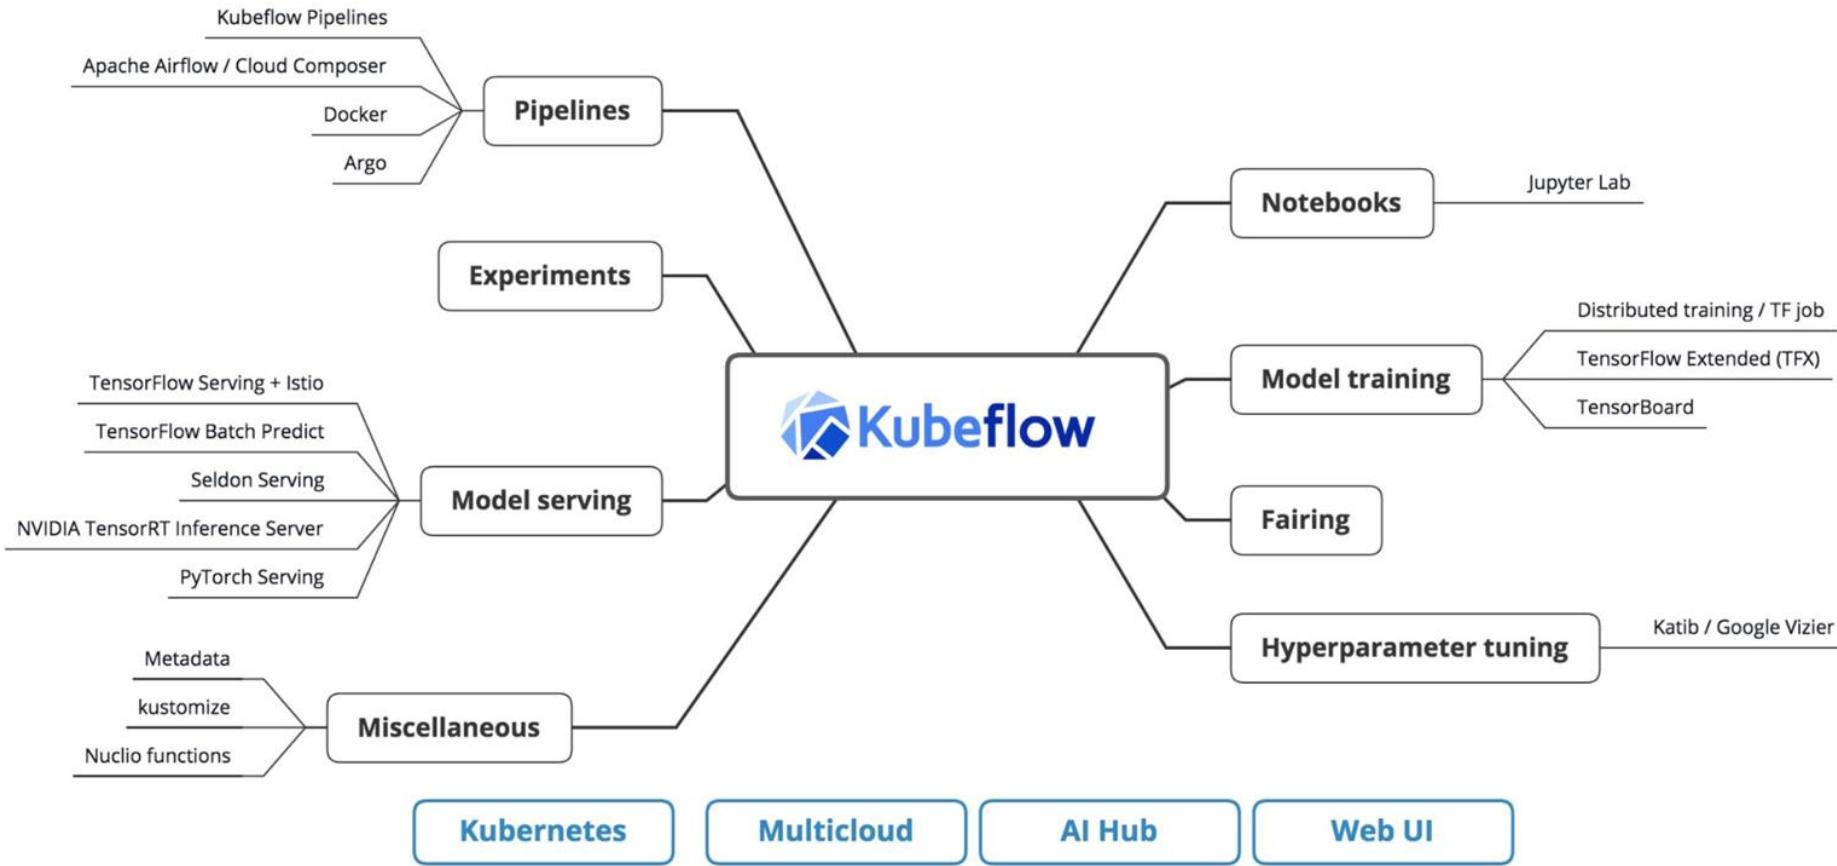
Open
design



Open
source



Open
to ideas



Kubeflow Mission Statement

**Makes it Easy for Everyone to Develop, Deploy and
Manage a Portable, Distributed and Scalable ML
system on Kubernetes**

KubeFlow 1.0 GA Release	
Stable Components	Beta Components
Central Dashboard, kfctl, Profile Controller, Notebooks, TFServing, Training Operators (Tensorflow, Pytorch), Docs	Pipelines, Metadata Store, HP Tuning, KFServing, Fairing, Kale, Training Operators(XGBoost, MxNet)

Code: Today

Local

```
import xgboost

class MyModel(object):
    def train(self):
        # load data
        # do feature engineering
        # train a model

    def predict():
        # prediction logic

if __name__ == '__main__':
    model = MyModel()
    model.train()
```

Build & Deploy to AI Platform

Training

```
gcloud ml-engine jobs submit training my_job \
    --module-name trainer.task \
    --staging-bucket gs://my-bucket \
    --package-path /my/code/path/trainer \
    --packages additional-dep1.tar.gz,dep2.whl
```

Prediction

```
gcloud alpha ml-engine versions create
{VERSION_NAME} --model {MODEL_NAME} \
--origin
gs://{BUCKET}/{MODEL_DIR}/ \
--runtime-version
{RUNTIME_VERSION} \
--package-uris
gs://{BUCKET}/{PACKAGES_DIR}/my_package-0.2.tar.gz \
--model-class=my_model.ModelExample
```

Build & Deploy to Kubeflow

```
apiVersion: kubeflow.org/v1alpha2
kind: TFJob
metadata:
  labels:
    experiment: experiment10
  name: tfjob
  namespace: kubeflow
spec:
  tfReplicaSpecs:
    Ps:
      replicas: 1
      template:
        metadata:
          creationTimestamp: null
        spec:
          containers:
            - args:
              - python
              - tf_cnn_benchmarks.py
            image:
            .
            .
            .
.
```

Code: With Kubeflow Fairing

Local

```
import xgboost

class MyModel(object):
    def train(self):
        # load data
        # train a model

    def predict():
        # prediction logic

from fairing import TrainJob
from fairing.backends import Backend

job = TrainJob(MyModel,
               backend=Backend("Local",
                               "fairing.config"))
job.submit()

endpoint = PredictionEndpoint(MyModel,
                               backend=Backend("Local",
                                               "fairing.config"))
endpoint.create()
```

Build & Deploy to AI Platform

```
import xgboost

class MyModel(object):
    def train(self):
        # load data
        # train a model

    def predict():
        # prediction logic

from fairing import TrainJob
from fairing.backends import Backend

job = TrainJob(MyModel,
               backend=Backend("ai_platform",
                               "fairing.config"))
job.submit()

endpoint = PredictionEndpoint(MyModel,
                               backend=Backend("ai_platform",
                                               "fairing.config"))
endpoint.create()
```

Build & Deploy to Kubeflow

```
import xgboost

class MyModel(object):
    def train(self):
        # load data
        # train a model

    def predict():
        # prediction logic

from fairing import TrainJob
from fairing.backends import Backend

job = TrainJob(MyModel,
               backend=Backend("Kubeflow",
                               "fairing.config"))
job.submit()

endpoint = PredictionEndpoint(MyModel,
                               backend=Backend("Kubeflow",
                                               "fairing.config"))
endpoint.create()
```

Kubeflow Fairing

An open source Hybrid
ML SDK for data
scientists to 'write ML
code once and run
anywhere'



Data scientist focused: Simple and uses language familiar Data Scientists



Multi-platform: Supports AI Platform and Kubeflow, making it easy for users to switch between on-prem and GCP



Scalable and cost-effective: Data Scientists can easily burst onto GCP when they need more resources (i.e. more machines, GPUs, or TPUs)

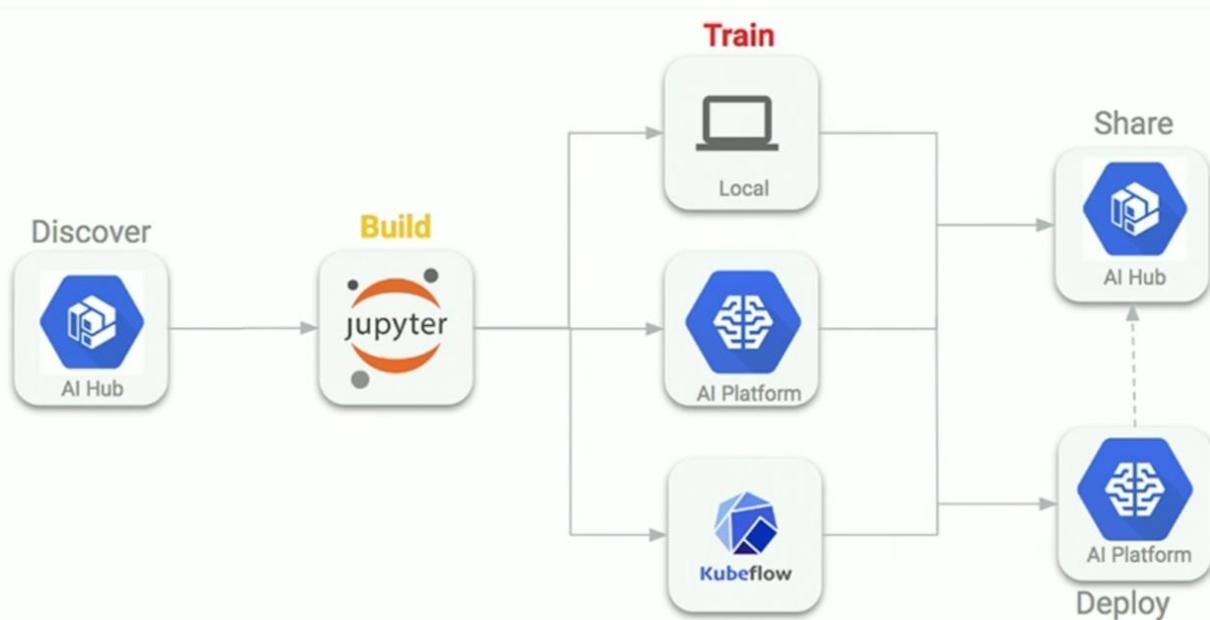


Easily train, tune, and deploy models:
Supports the full ML lifecycle



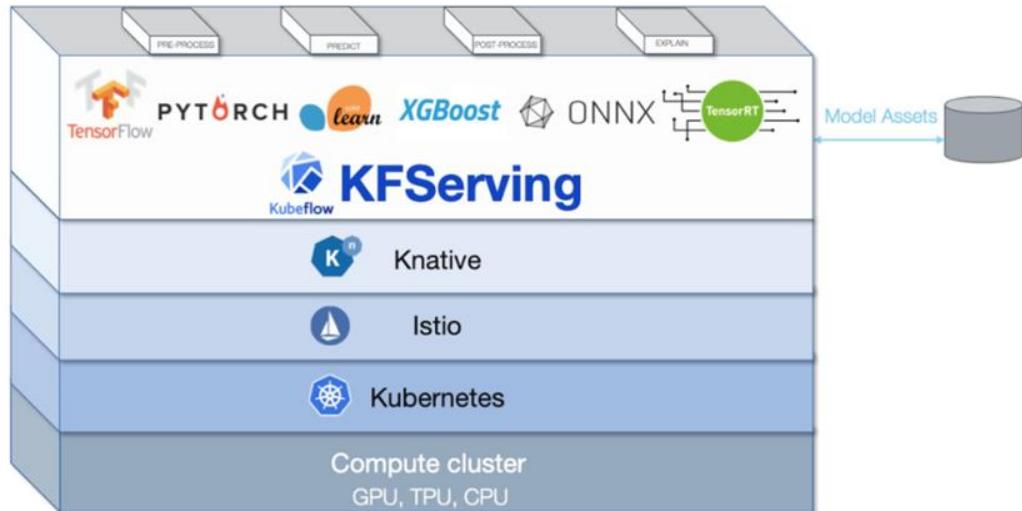
Multi-framework: Supports XGBoost, TensorFlow (single node) and Pytorch (single node)

Sample: Hybrid E2E ML Kubeflow Fairing



KFServing

- Scalable,
Kubernetes-native
intererencing
- Usage pattern:
 - High Availability
 - Quick addition of capacity
 - Potentially need GPUs



A Clean Interface

```
apiVersion: "serving.kubeflow.org/v1alpha1"
kind: "InferenceService"
metadata:
  name: "sklearn-iris"
spec:
  default:
    sklearn:
      storageUri: "gs://kfserving-samples/models/sklearn/iris"
```

```
apiVersion: "serving.kubeflow.org/v1alpha1"
kind: "InferenceService"
metadata:
  name: "flowers-sample"
spec:
  default:
    tensorflow:
      storageUri: "gs://kfserving-samples/models/tensorflow/flowers"
```

```
apiVersion: "serving.kubeflow.org/v1alpha1"
kind: "KFService"
metadata:
  name: "pytorch-cifar10"
spec:
  default:
    pytorch:
      storageUri: "gs://kfserving-samples/models/pytorch/cifar10"
      modelClassName: "Net"
```





1. Containerise
2. Deploy
3. Monitor



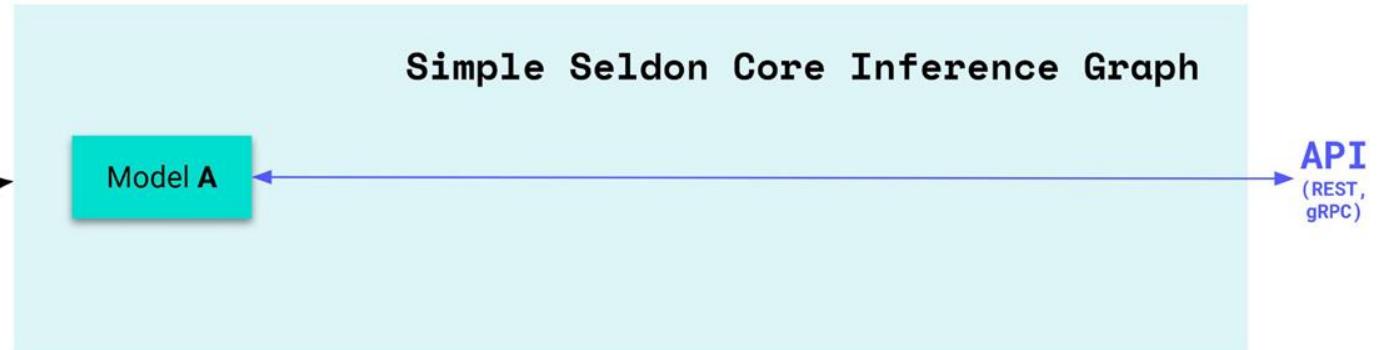
From model binary

Or language wrapper

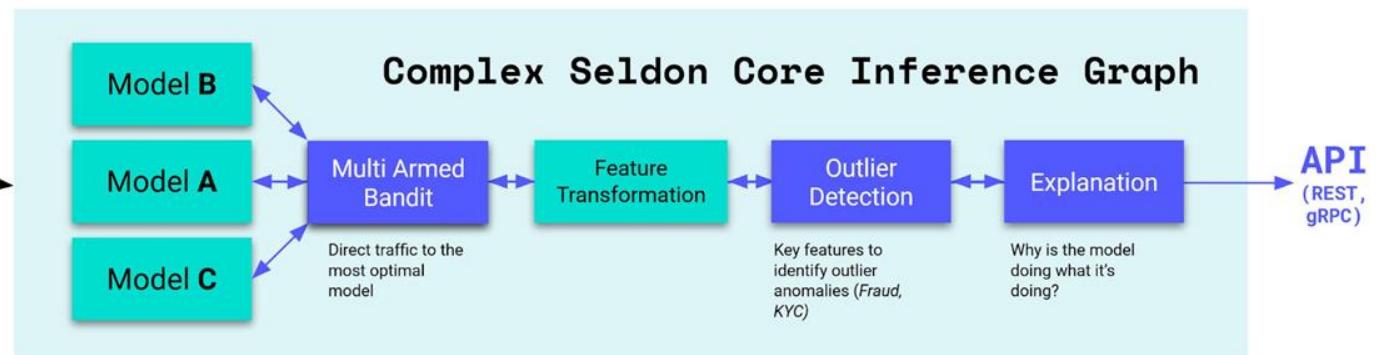


Into fully fledged microservice

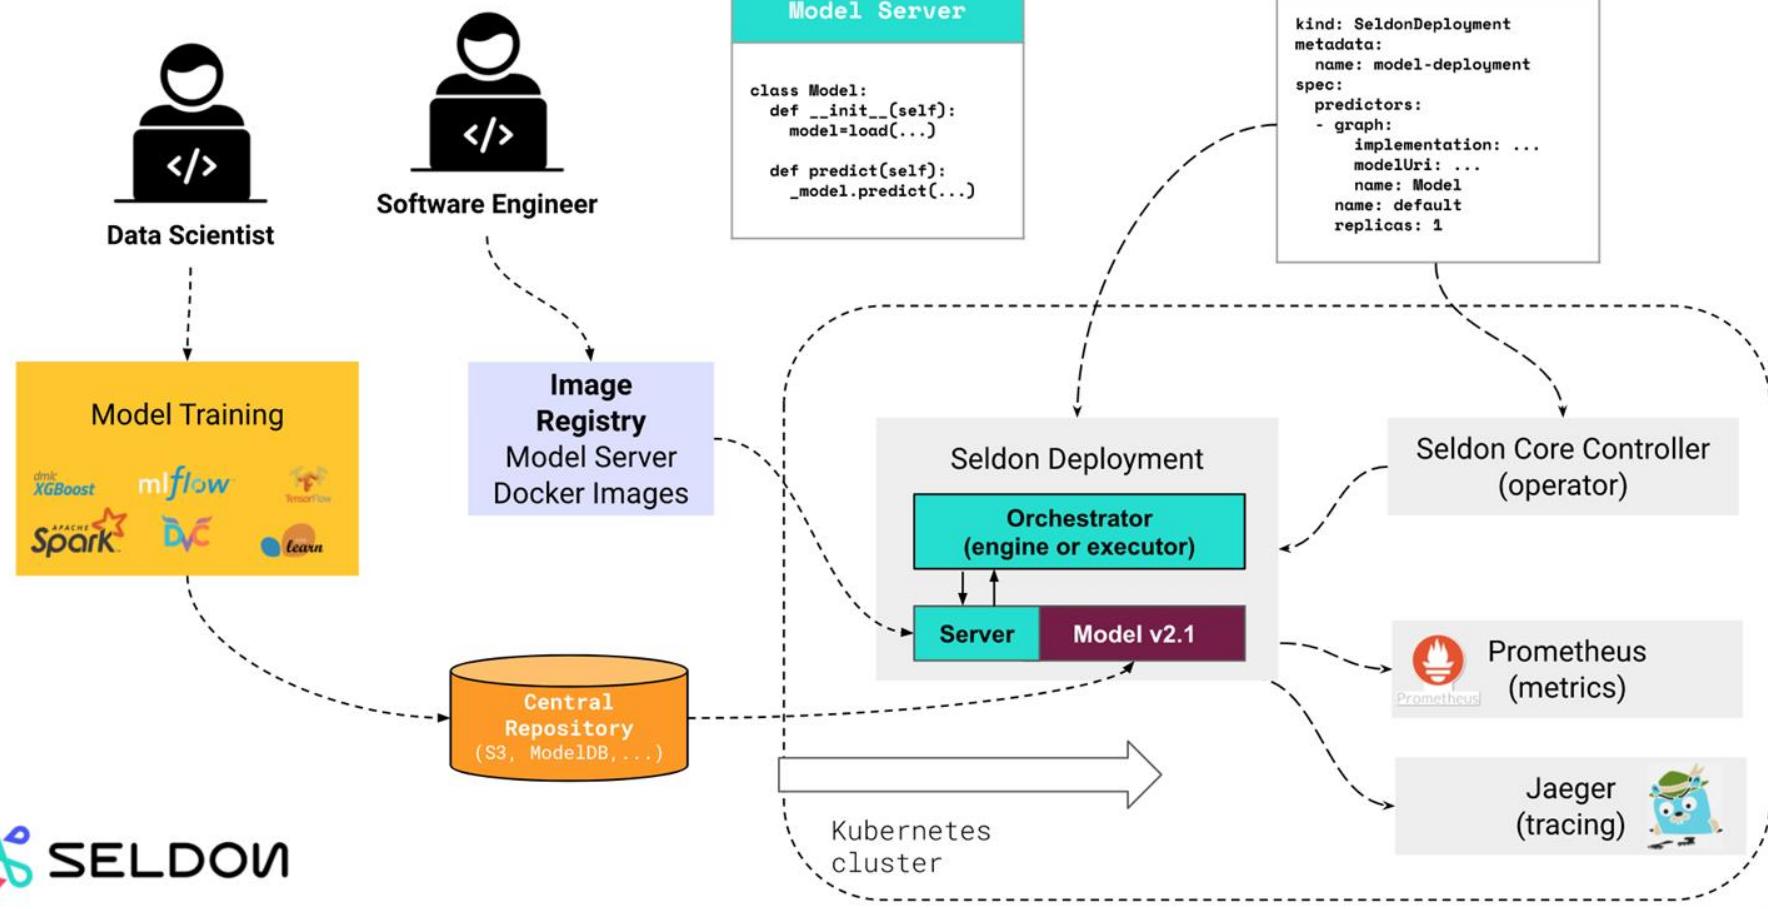
Simple Seldon Core Inference Graph



Complex Seldon Core Inference Graph



E2E Model Serving





Production ML Pipelines in the Cloud

TensorFlow Extended



What we're doing

- We'll build a production-style end-to-end machine learning pipeline for training models
- We'll build it with TensorFlow Extended (TFX) on Google Cloud AI Platform



Very High Level Architecture

TensorFlow

TensorFlow Extended (TFX)

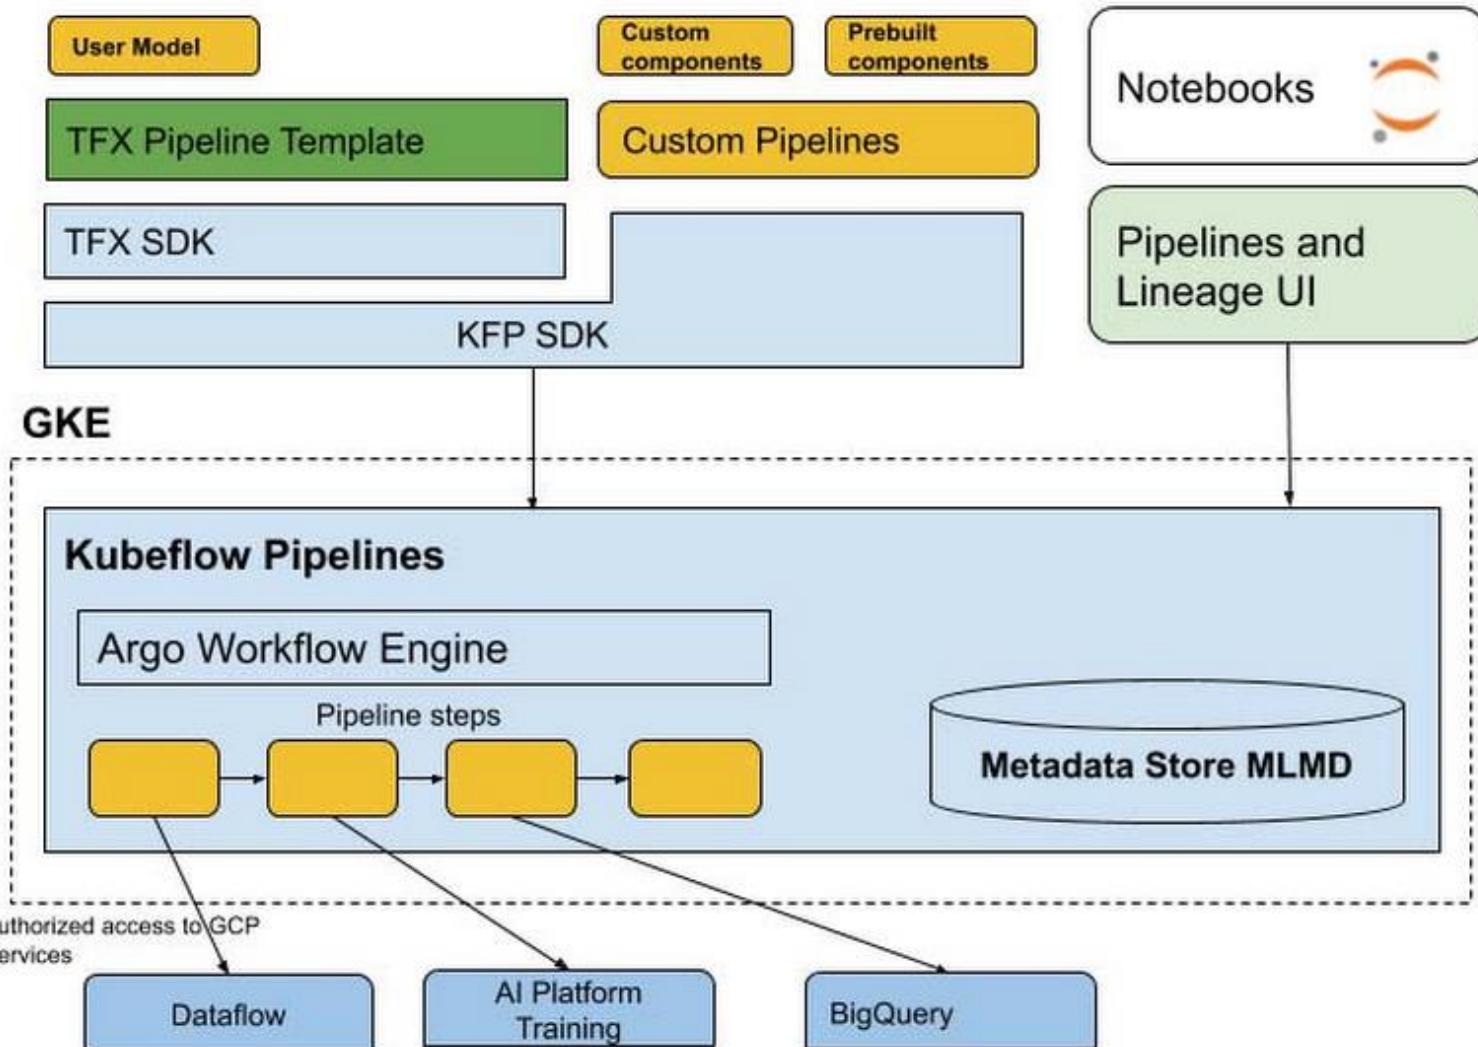
Kubeflow Pipelines

Google Kubernetes Engine (GKE)

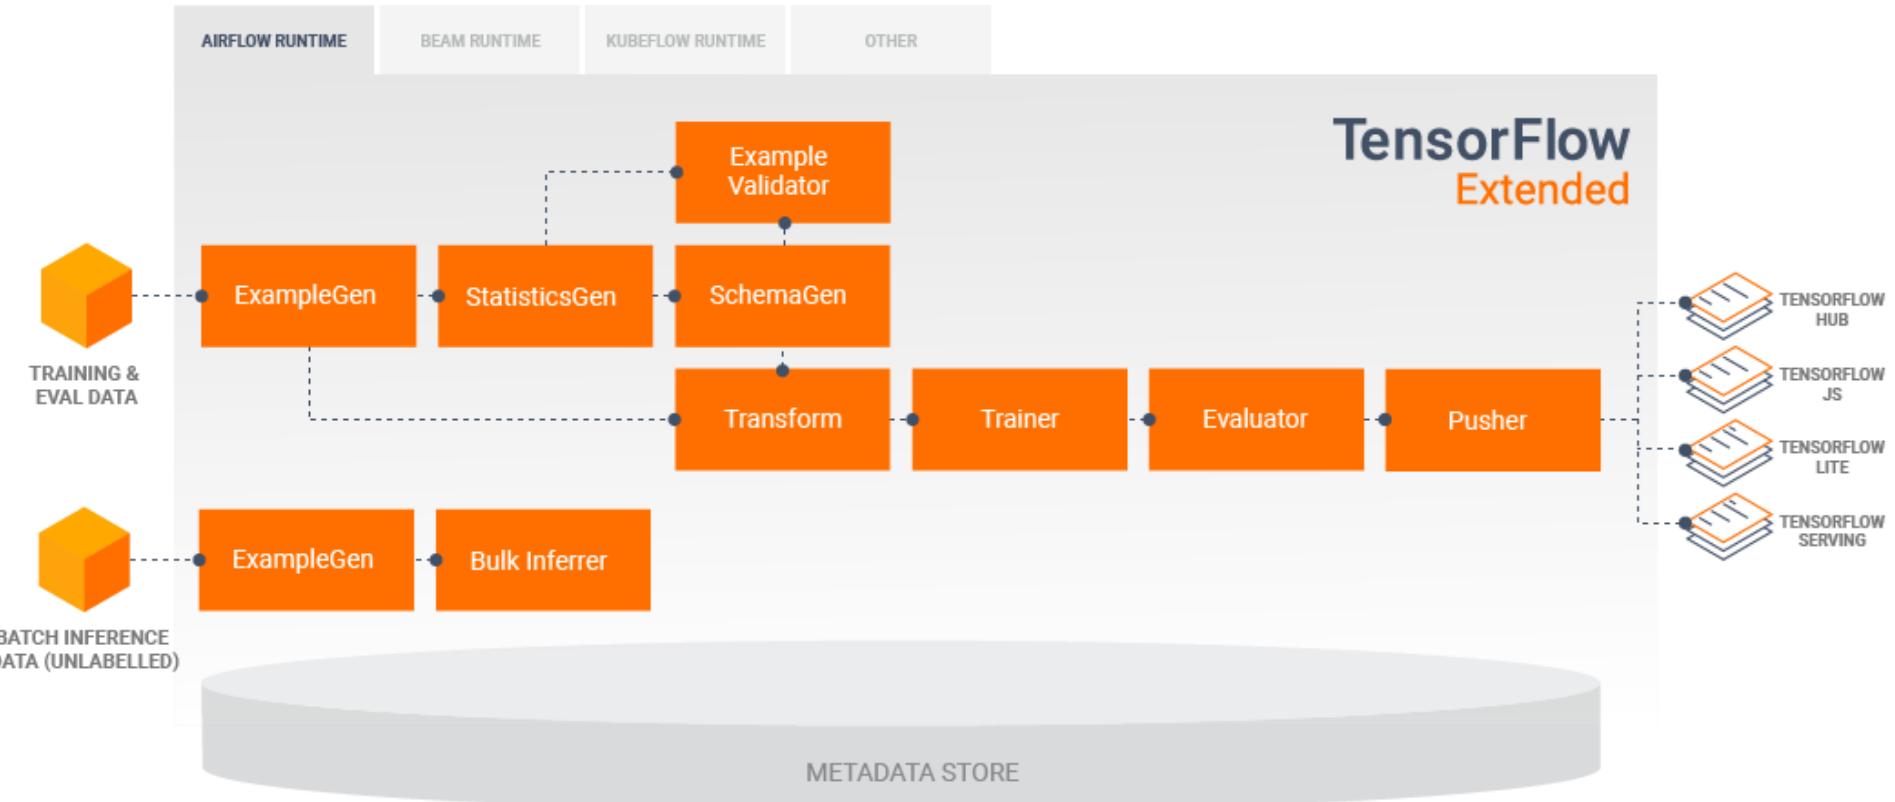
GCS

BigQuery

Dataflow



TFX CONFIG





Kubeflow Pipelines



Kubeflow Pipelines

A platform for building and deploying portable, scalable machine learning (ML) workflows based on Docker containers on Kubernetes.

Python SDK	Kubernetes runtime stack	Interactive UI
User-facing library for defining components and pipelines	Control plane server with Kubernetes-native workflow orchestration engine based on Argo, including custom extension for scheduled (cron-like) workflows.	Custom UI for launching & viewing pipeline runs, experiments, comparing output artifacts with out-of-the-box visualization capabilities.



TFX End-to-End Example

Chicago Taxi Cab Dataset



TFX End-to-End Example

Chicago Taxi Cab Dataset

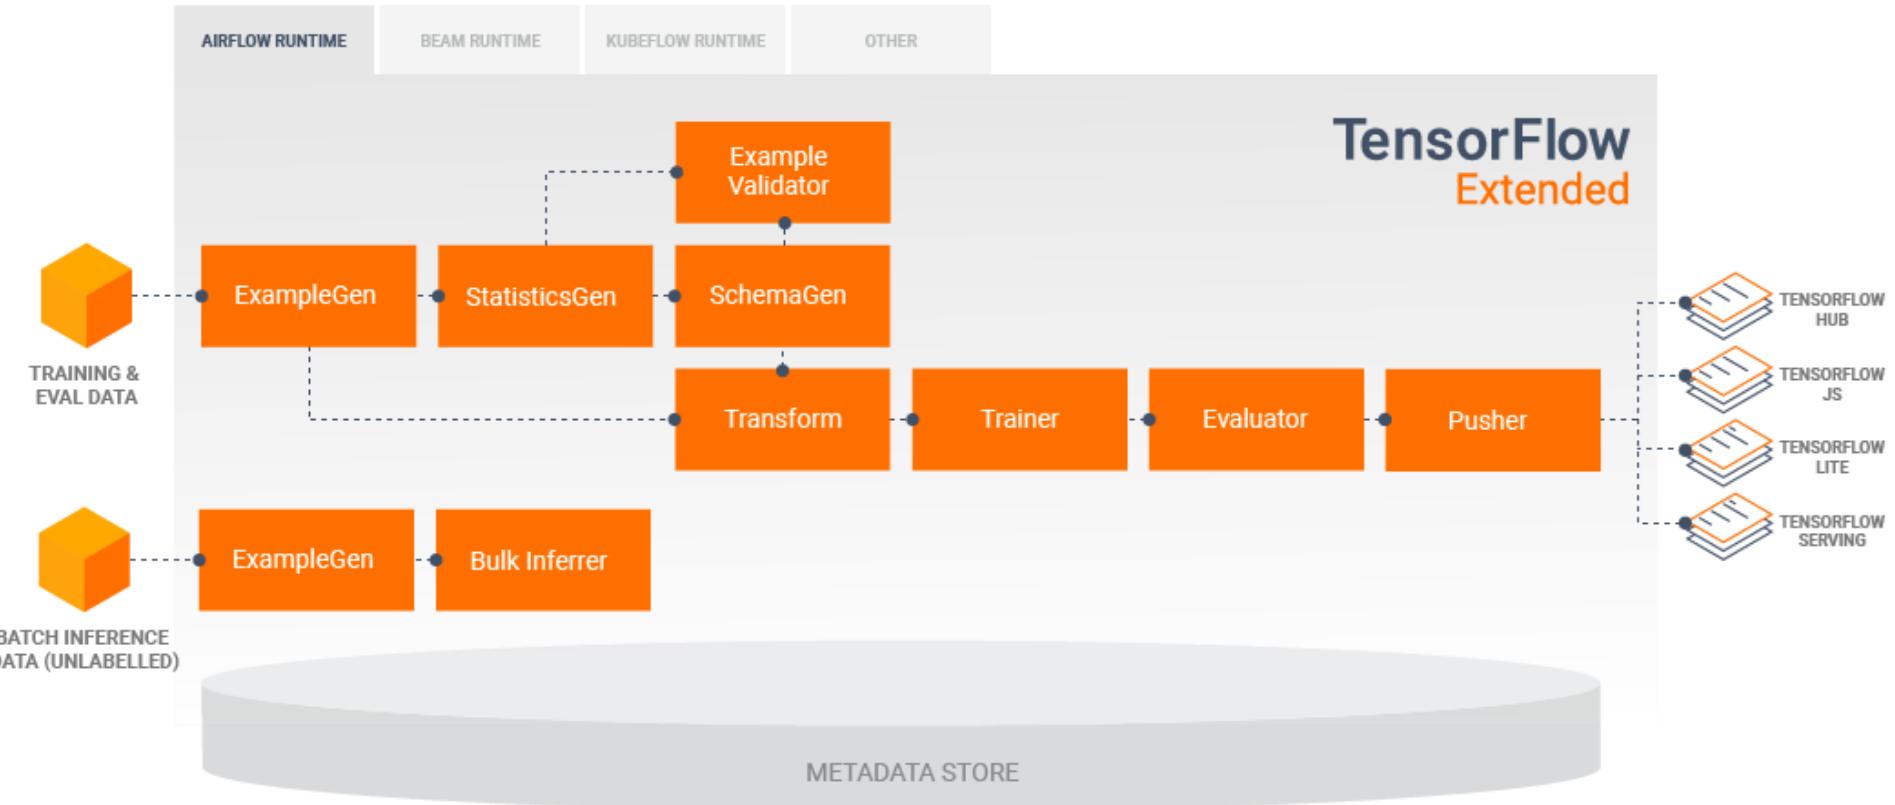


Features

Categorical Features	Bucket Features	Vocab Features	Dense Float Features
trip_start_hour	pickup_latitude	payment_type	trip_miles
trip_start_day	pickup_longitude	company	fare
trip_start_month	dropoff_latitude		trip_seconds
pickup/dropoff_census_tract	dropoff_longitude		
pickup/dropoff_community_area			

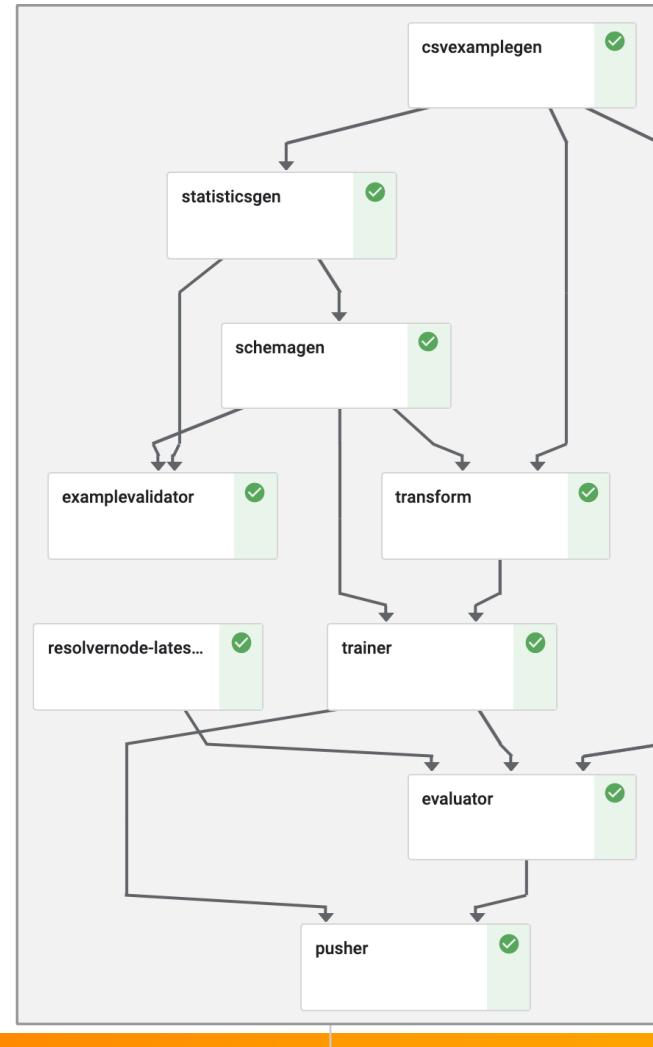
Label = tips > (fare * 20%)

TFX CONFIG





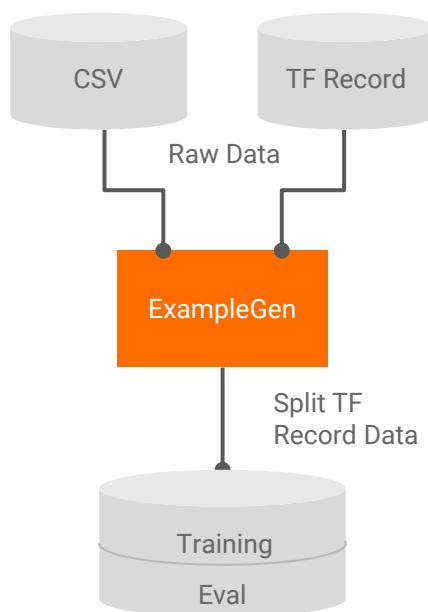
TFX End-to-End Example





Component: ExampleGen

Inputs and Outputs



Configuration

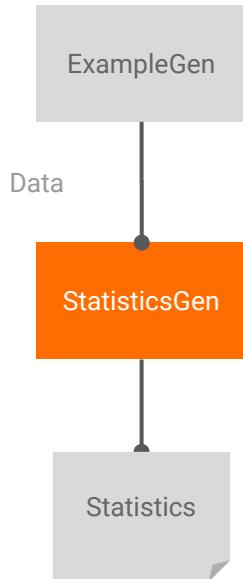
```
examples = csv_input(os.path.join(data_root, 'simple'))  
example_gen = CsvExampleGen(input_base=examples)
```





Component: StatisticsGen

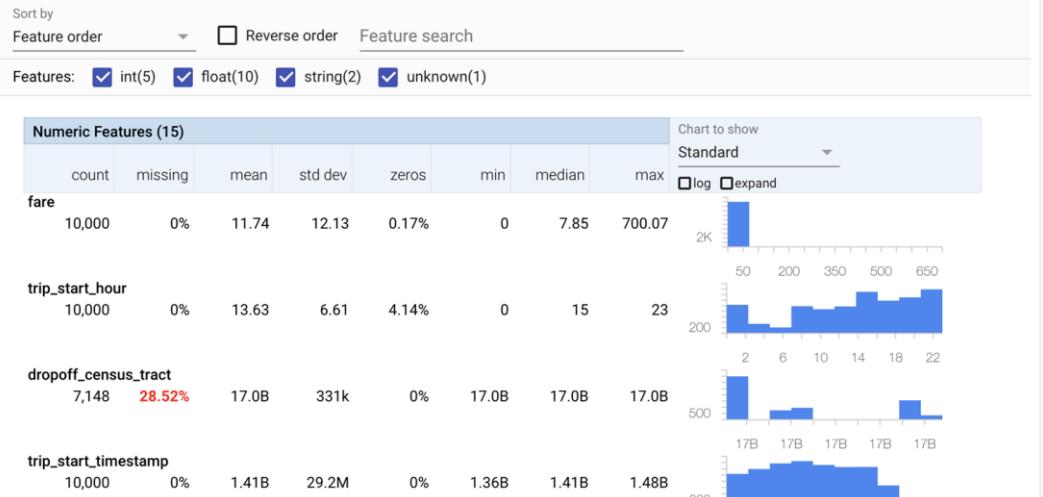
Inputs and Outputs



Configuration

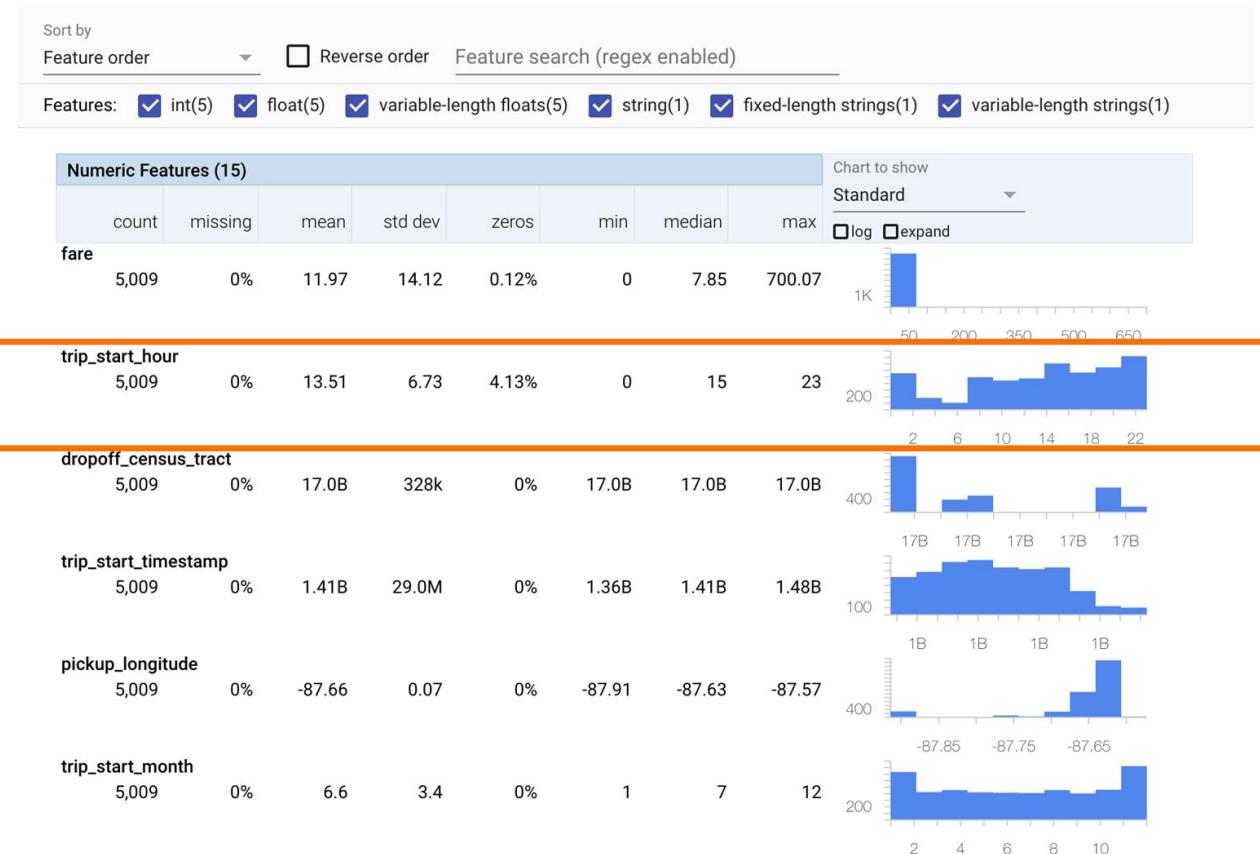
```
statistics_gen =  
    StatisticsGen(input_data=example_gen.outputs.examples)
```

Visualization





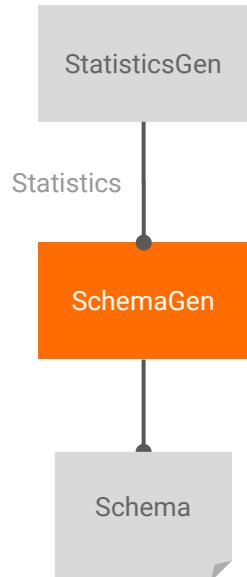
Analyzing Data with TensorFlow Data Validation





Component: SchemaGen

Inputs and Outputs



Configuration

```
infer_schema = SchemaGen(stats=statistics_gen.outputs.output)
```

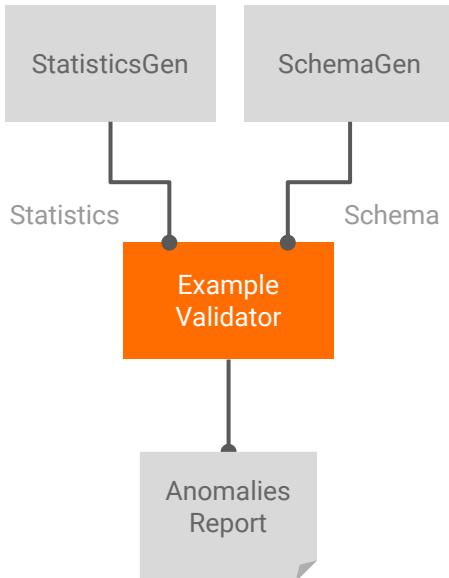
Visualization

Feature name	Type	Presence	Valency	Domain
'fare'	FLOAT	required	single	-
'trip_start_hour'	INT	required	single	-
'pickup_census_tract'	BYTES	optional		-
'dropoff_census_tract'	FLOAT	optional	single	-
'company'	STRING	optional	single	'company'



Component: ExampleValidator

Inputs and Outputs



Configuration

```
validate_stats = ExampleValidator(  
    stats=statistics_gen.outputs.output,  
    schema=infer_schema.outputs.output)
```

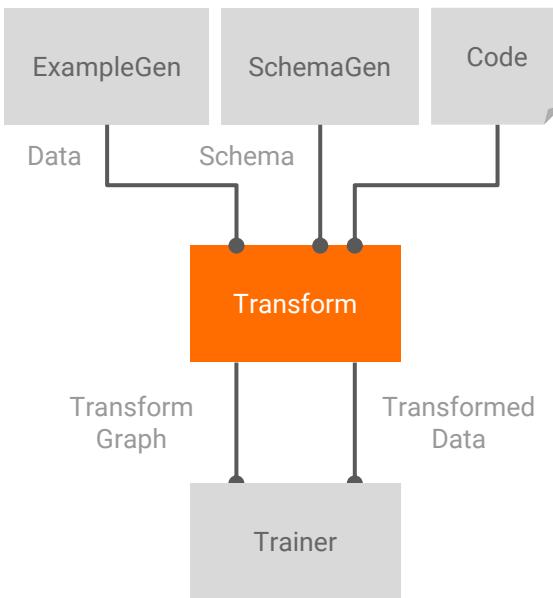
Visualization

Feature name	Anomaly short description	Anomaly long description
'payment_type'	Unexpected string values	Examples contain values missing from the schema: Prcard (<1%).
'company'	Unexpected string values	Examples contain values missing from the schema: 2092 - 61288 Sbeih company (<1%), 2192 - 73487 Zeymane Corp (<1%), 2192 - Zeymane Corp (<1%), 2823 - 73307 Seung Lee (<1%), 3094 - 24059 G.L.B. Cab Co (<1%), 3319 - CD Cab Co (<1%), 3385 - Erman Cab (<1%), 3897 - 57856 Ilie Malec (<1%), 4053 - 40193 Adwar H. Nikola (<1%), 4197 - Royal Star (<1%), 585 - 88805 Valley Cab Co (<1%), 5874 - Sergey Cab Corp. (<1%), 6057 - 24657 Richard Addo (<1%), 6574 - Babylon Express Inc. (<1%), 6742 - 83735 Tasha ride inc (<1%).



Component: Transform

Inputs and Outputs



Configuration

```
transform = Transform(  
    input_data=example_gen.outputs.examples,  
    schema=infer_schema.outputs.output,  
    module_file=taxi_module_file)
```

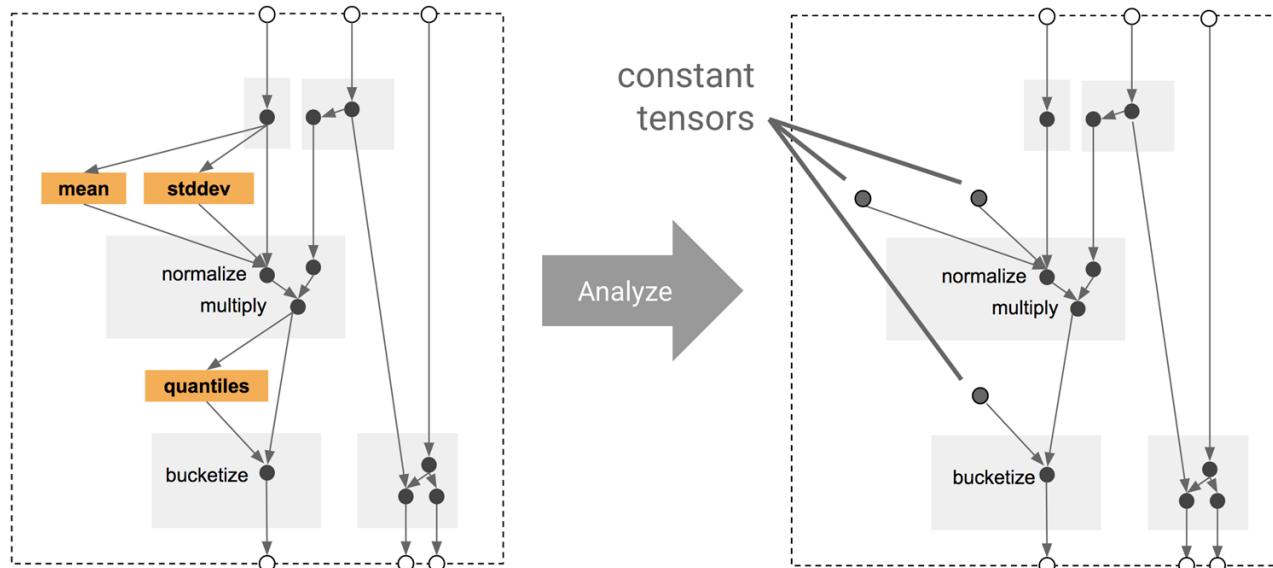
Code

```
for key in _DENSE_FLOAT_FEATURE_KEYS:  
    outputs[_transformed_name(key)] = transform.scale_to_z_score(  
        _fill_in_missing(inputs[key]))  
    # ...  
  
outputs[_transformed_name(_LABEL_KEY)] = tf.where(  
    tf.is_nan(taxi_fare),  
    tf.cast(tf.zeros_like(taxi_fare), tf.int64),  
    # Test if the tip was > 20% of the fare.  
    tf.cast(  
        tf.greater(tips, tf.multiply(taxi_fare, tf.constant(0.2))), tf.int64))  
    # ...
```



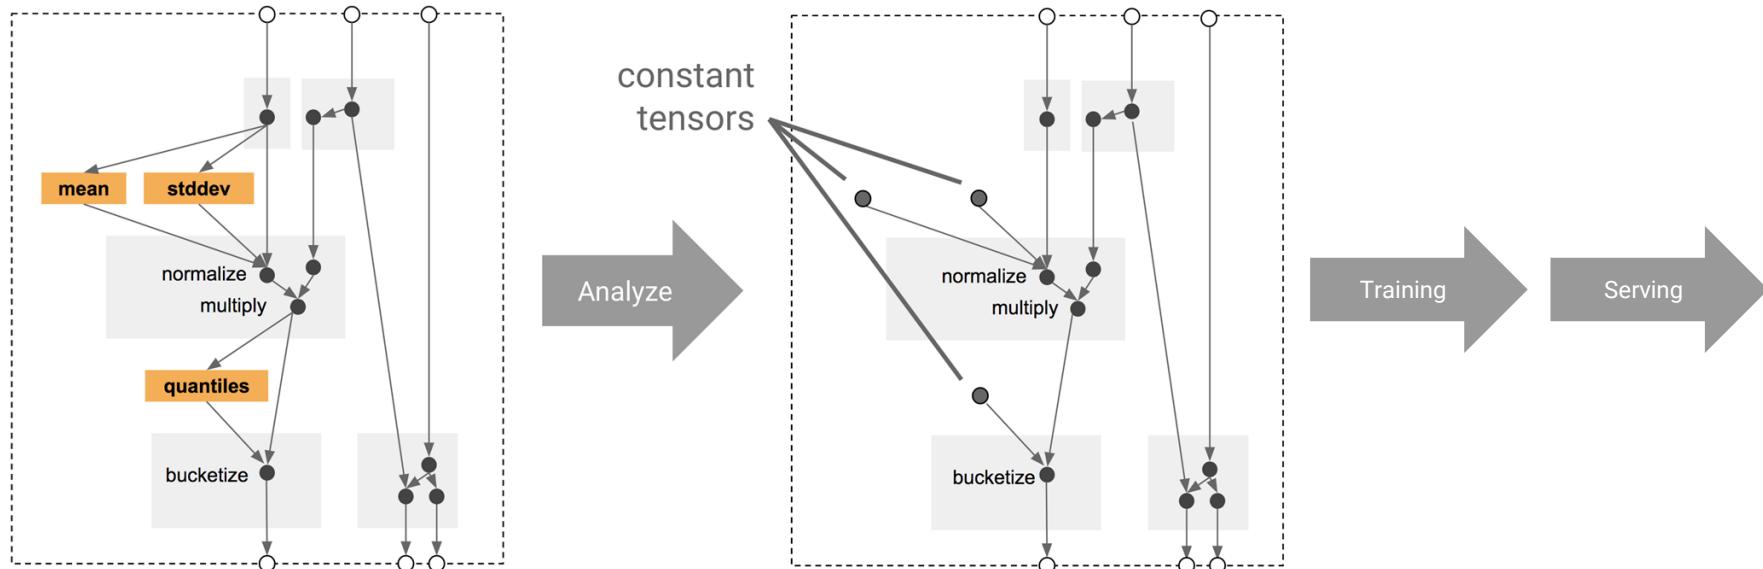


Using TensorFlow Transform for Feature Engineering





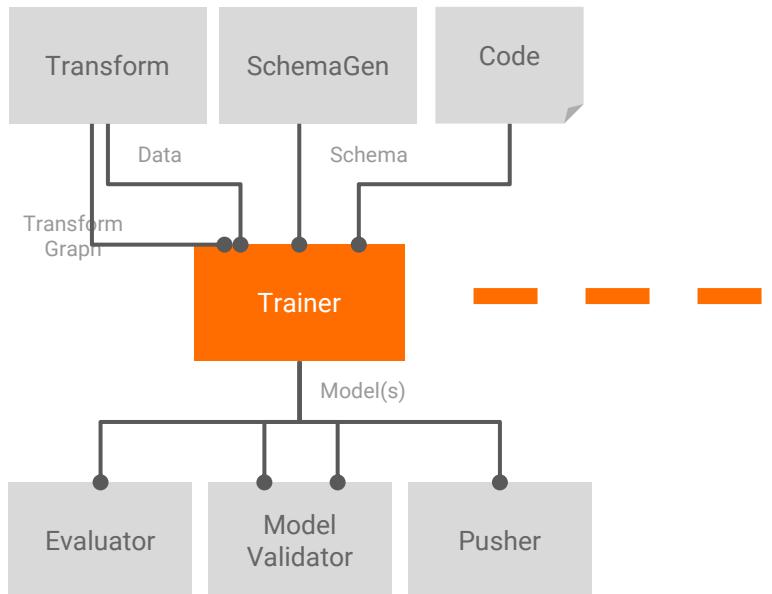
Using TensorFlow Transform for Feature Engineering





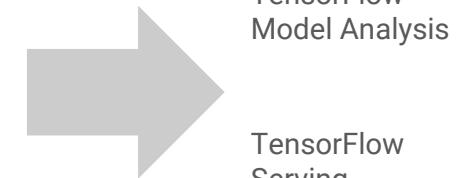
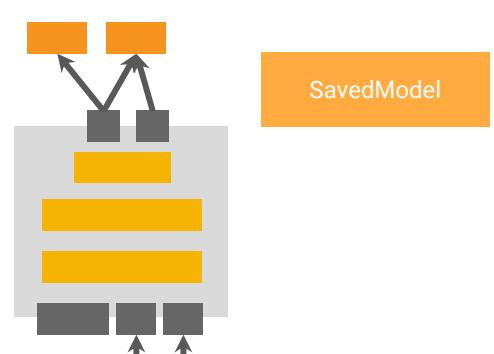
Component: Trainer

Inputs and Outputs



Highlight: SavedModel Format

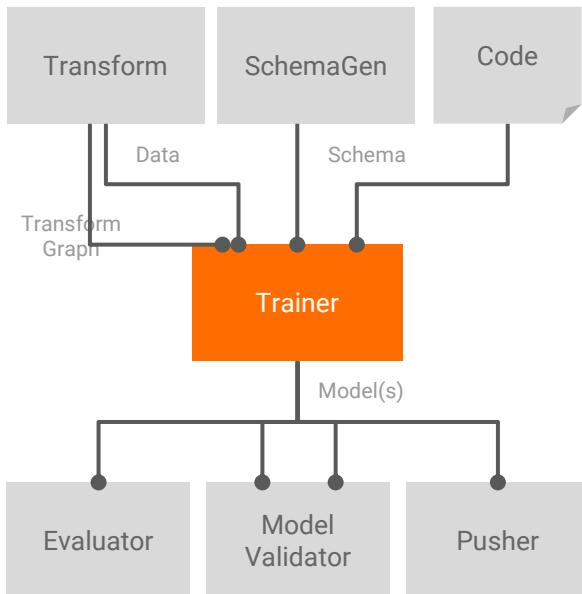
Train, Eval, and Inference Graphs





Component: Trainer

Inputs and Outputs



Configuration

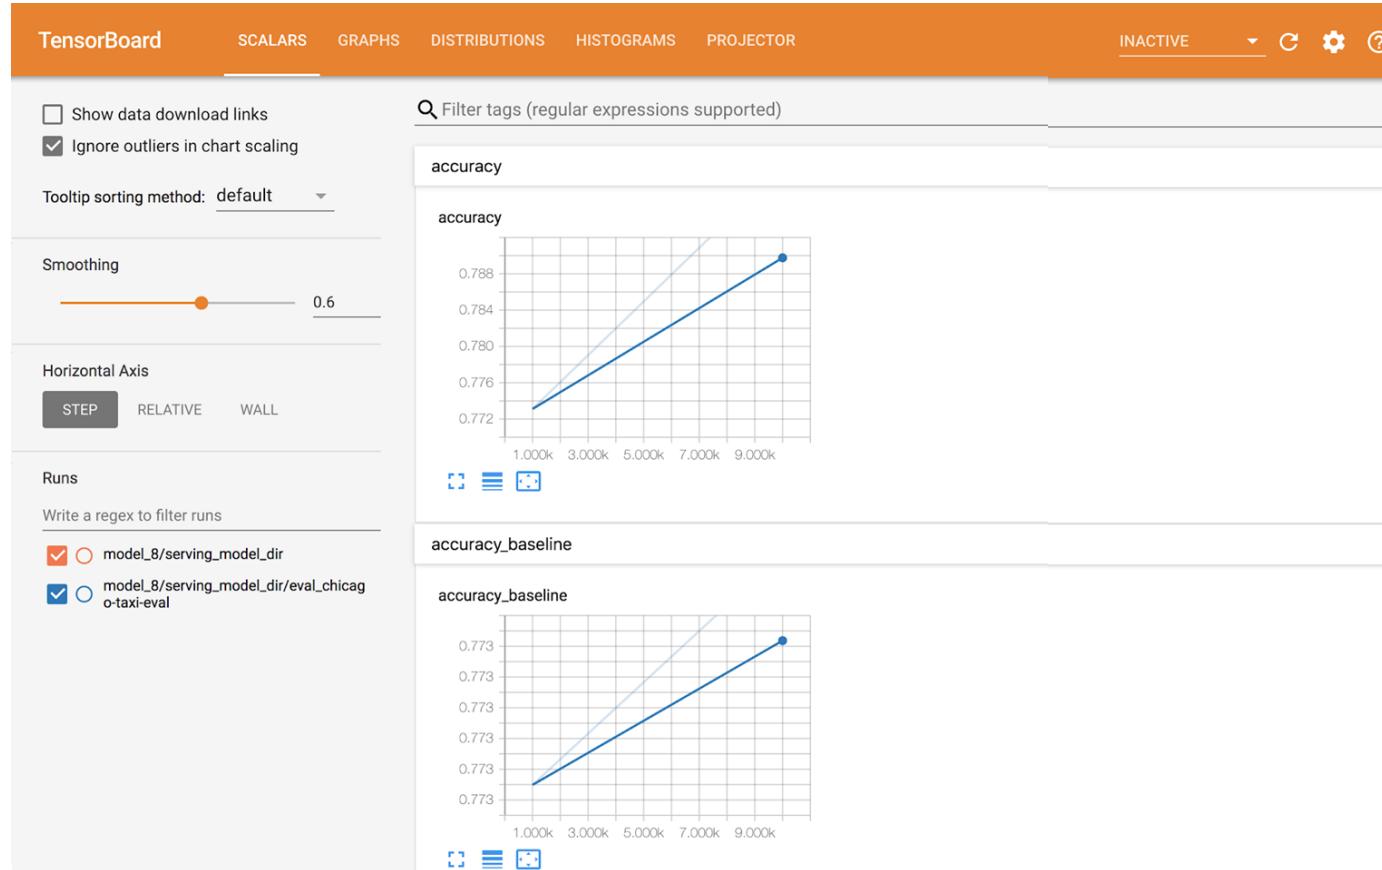
```
trainer = Trainer(  
    module_file=taxi_module_file,  
    transformed_examples=transform.outputs.transformed_examples,  
    schema=infer_schema.outputs.output,  
    transform_output=transform.outputs.transform_output,  
    train_steps=10000,  
    eval_steps=5000,  
    warm_starting=True)
```

Code

Just TensorFlow :)



```
# Open up Tensorboard for model_id.  
print(display_tensorboard(model_id))  
  
http://your.host.name:53143
```





```
# Compare Tensorboard metrics for different models.  
if num_models > 1:  
    print(display_tensorboard(model_id, other_model_id=other_model_id))  
  
http://your.host.name:53230
```

TensorBoard SCALARS GRAPHS DISTRIBUTIONS HISTOGRAMS PROJECTOR INACTIVE ▾ C G ?

Show data download links
 Ignore outliers in chart scaling

Tooltip sorting method: default ▾

Smoothing: 0.6

Horizontal Axis: STEP RELATIVE WALL

Runs: Write a regex to filter runs

model_8/serving_model_dir
 model_8/serving_model_dir/eval_chicag
 model_20/serving_model_dir
 model_20/serving_model_dir/eval_chica

accuracy

accuracy

Step	Run 1 Accuracy	Run 2 Accuracy
1.000k	~0.840	~0.770
3.000k	~0.850	~0.775
5.000k	~0.860	~0.780
7.000k	~0.870	~0.785
9.000k	~0.880	~0.790

accuracy_baseline

accuracy_baseline

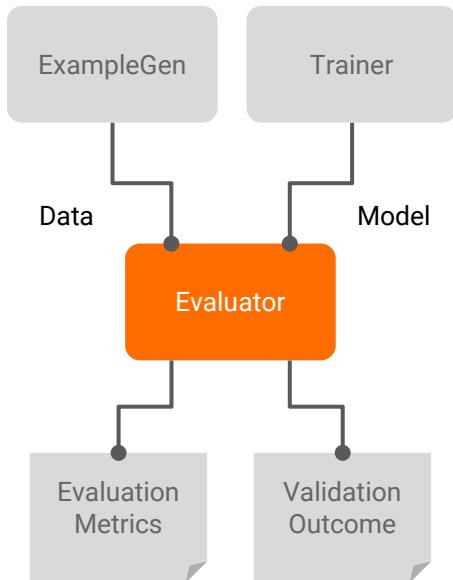
Step	Run 1 Accuracy	Run 2 Accuracy
1.000k	~0.784	~0.773
3.000k	~0.784	~0.773
5.000k	~0.784	~0.773
7.000k	~0.784	~0.773
9.000k	~0.784	~0.773

TOGGLE ALL RUNS



Component: Evaluator

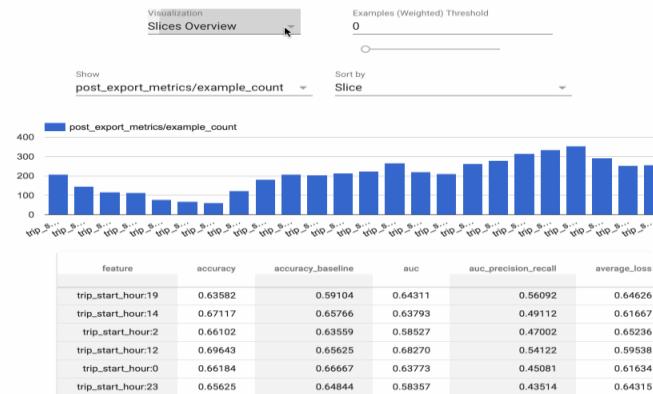
Inputs and Outputs



Configuration

```
evaluator = Evaluator(  
    examples=example_gen.outputs['examples'],  
    model=trainer.outputs['model'],  
    baseline_model=model_resolver.outputs['model'],  
    eval_config=eval_config)
```

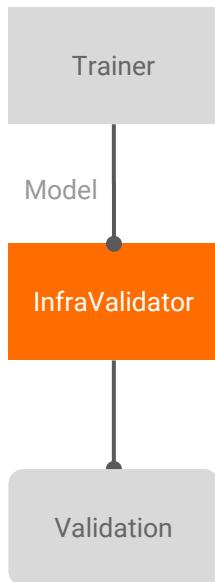
Visualization





Component: InfraValidator

Inputs and Outputs



Configuration

```
infra_validator = InfraValidator(  
    model=trainer.outputs['model'],  
    serving_spec=ServingSpec(...),  
    validation_spec=ValidationSpec(...),  
    request_spec=RequestSpec(...)  
)
```

Configuration Options

ServingSpec: Type of model server and infrastructure to test with

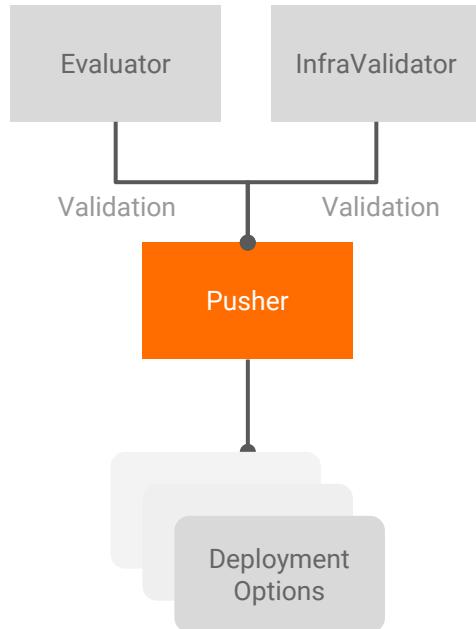
ValidationSpec: Adjusts the infra validation criteria or workflow

RequestSpec: Which model signature, how many examples to test



Component: Pusher

Inputs and Outputs



Configuration

```
pusher = Pusher(  
    model=trainer.outputs['model'],  
    model_blessing=evaluator.outputs['blessing'],  
    infra_blessing=infra_validator.outputs['blessing'],  
    push_destination=pusher_pb2.PushDestination(  
        filesystem=pusher_pb2.PushDestination.Filesystem(  
            base_directory=_serving_model_dir)))
```

Block push on validation outcome

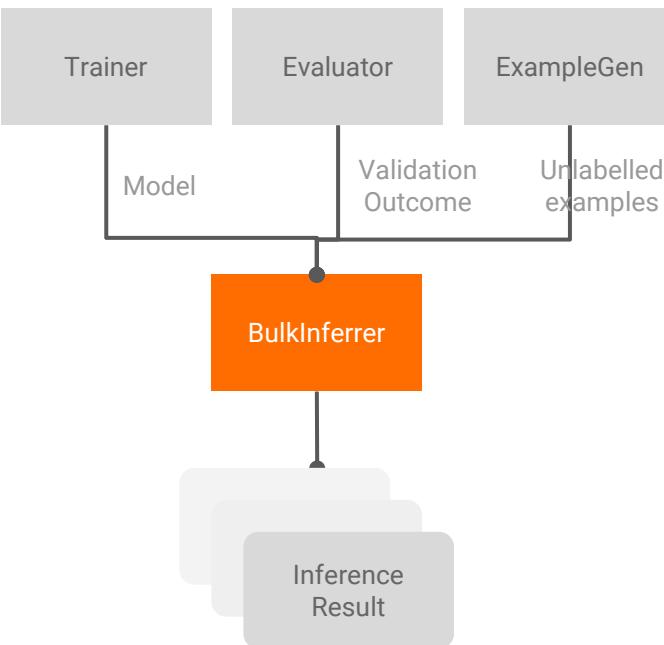
Push destinations supported today

- Filesystem (TensorFlow Lite, TensorFlow JS)
- TensorFlow Serving



Component: BulkInferencer

Inputs and Outputs



Configuration

```
bulk_inferrer = BulkInferencer(  
    examples=inference_example_gen.outputs['examples'],  
    model_export=trainer.outputs['output'],  
    model_blessing=evaluator.outputs['blessing'],  
    data_spec=bulk_inferrer_pb2.DataSpec(  
        example_splits=['unlabelled']),  
    model_spec=bulk_inferrer_pb2.ModelSpec())
```

Configuration Options

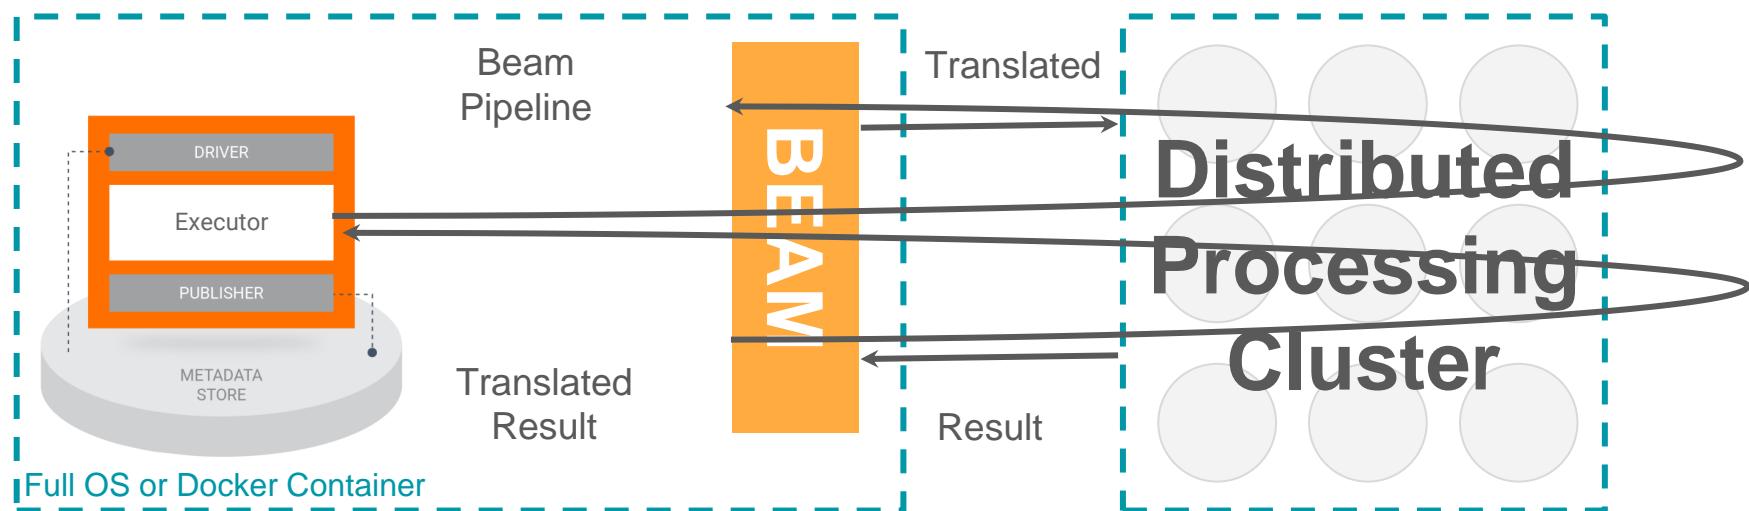
- Block batch inference on a successful model validation.
- Choose the inference examples from example gen's output.
- Choose the signatures and tags of inference model.

Inference Result

Contains features and predictions.



How TFX Components Use Beam



TFX Pipeline Nodes



What are Pipeline Nodes?

Pipeline nodes are special-purpose classes for performing advanced metadata operations

- Import external artifacts into ML-Metadata
- Perform queries of current ML Metadata based on artifact properties and history



ImporterNode

ImporterNode registers an external resource into MLMD so that downstream nodes can use the registered artifact as input

- Key requirements: source_uri, artifact_type

```
importer = ImporterNode(  
    instance_name='import_schema',  
    source_uri='uri/to/schema'  
    artifact_type=standard_artifacts.Schema,  
    reimport=False)
```



ResolverNode

ResolverNode is a special TFX node which components use to perform artifact resolution queries

- Key requirements: Class, Query config, Channel

```
latest_five_examples_resolver = ResolverNode(  
    instance_name='latest_five_examples_resolver',  
    resolver_class=latest_artifacts_resolver.LatestArtifactsResolver, # Class  
    resolver_config={'desired_num_of_artifacts' : 5}, # Query config  
    examples=example_gen.outputs['examples']) # Channel
```

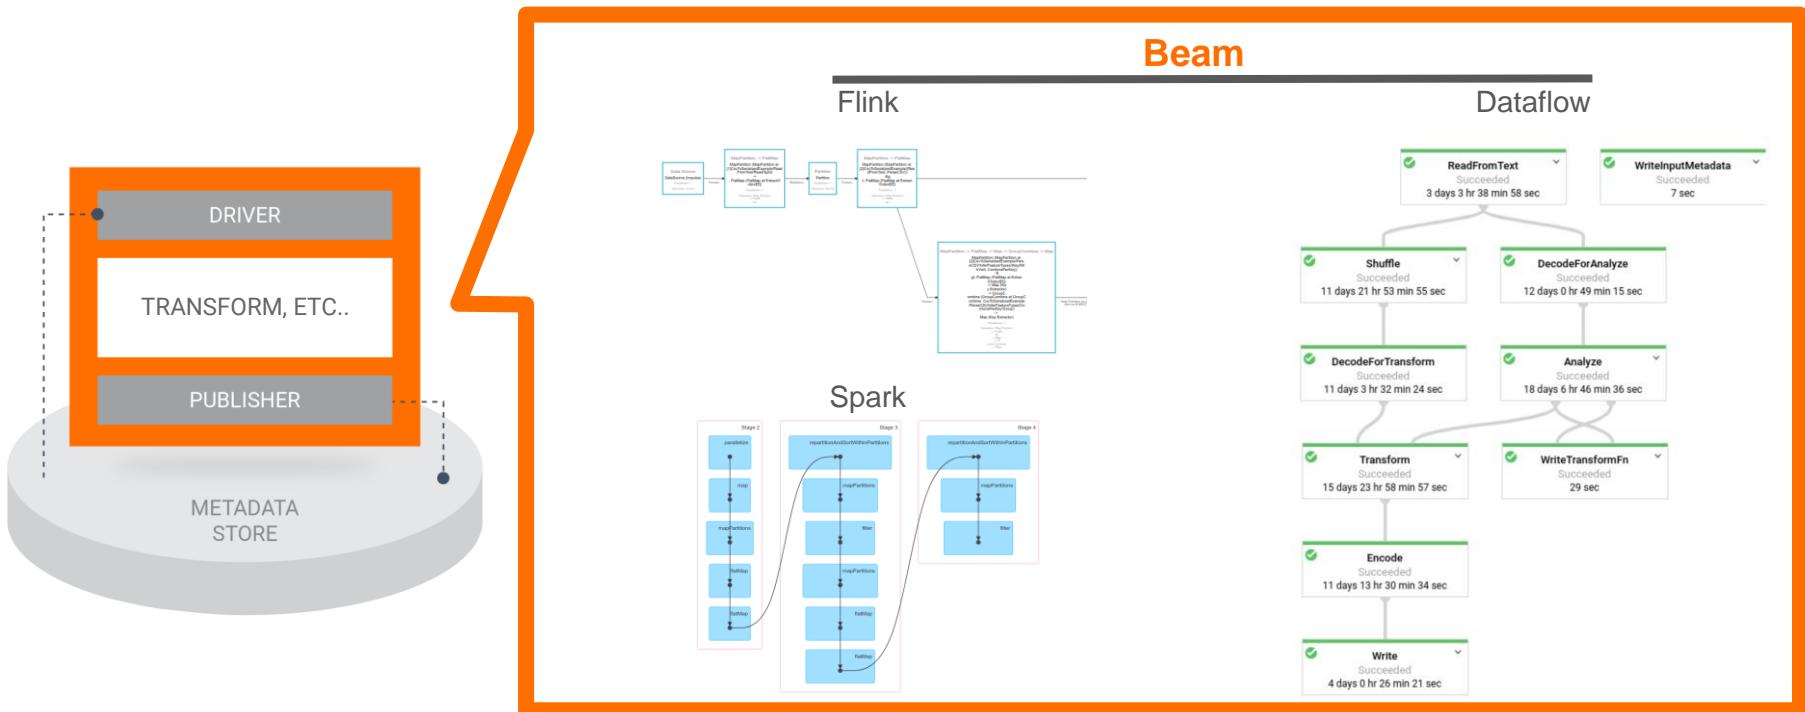


Current ResolverNodes

- **LatestArtifactsResolver**
 - Resolver that returns the latest n artifacts in a given channel
- **LatestBlessedModelResolver**
 - Returns the latest blessed model

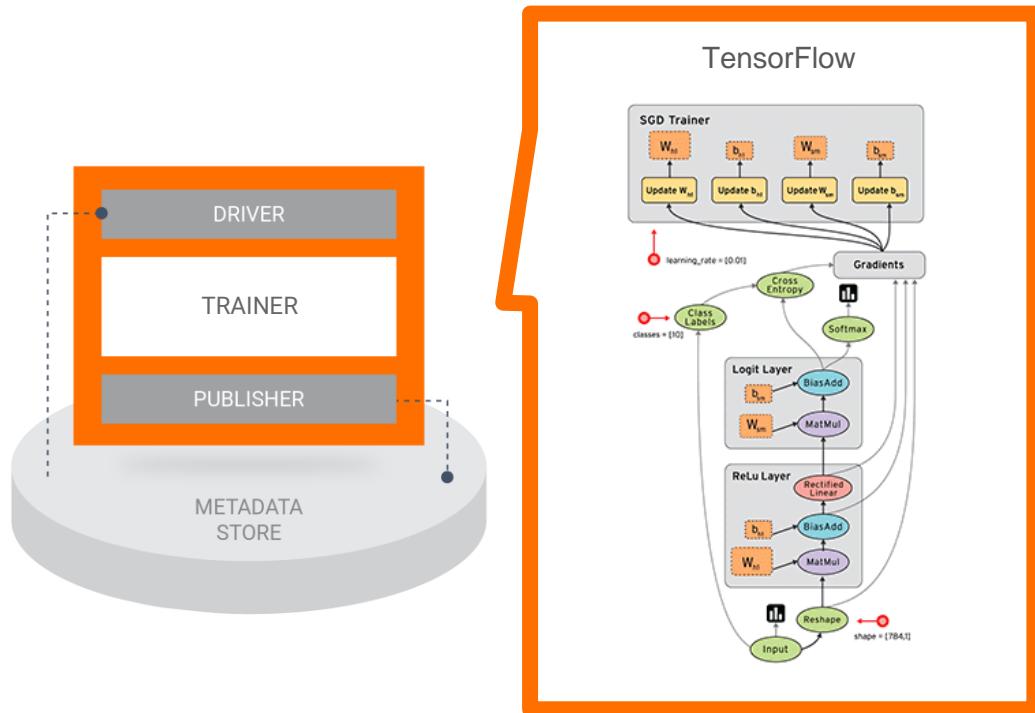


Executors do the work



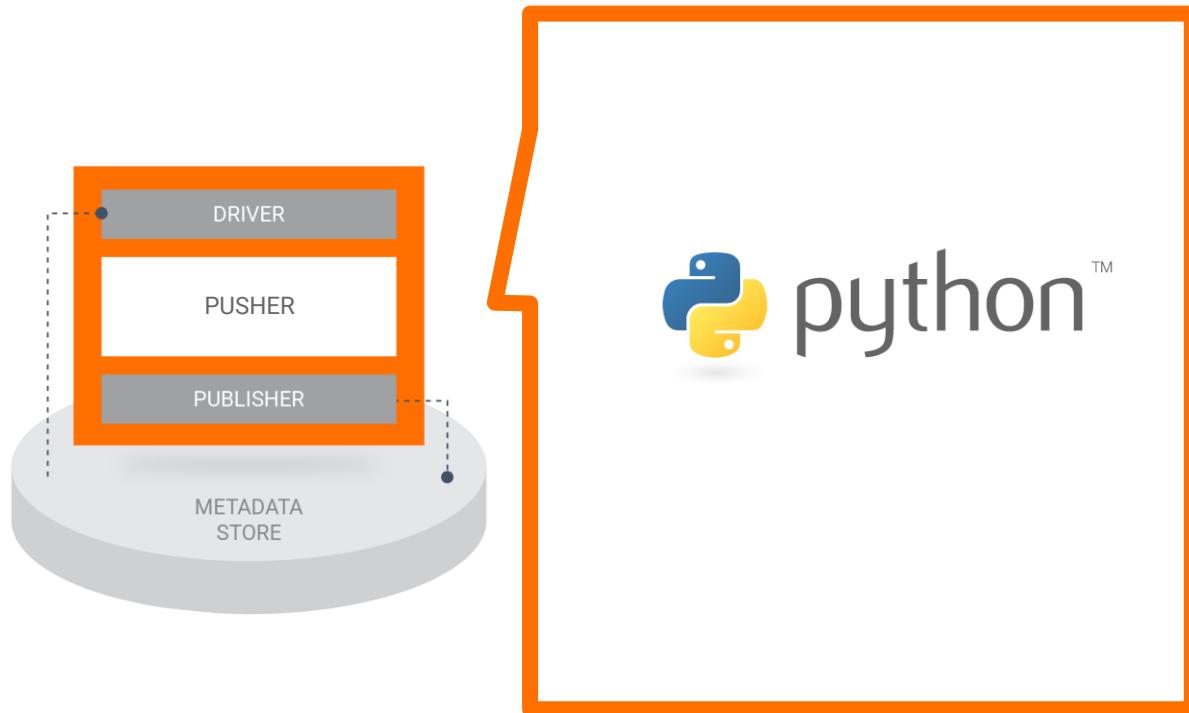


Executors do the work





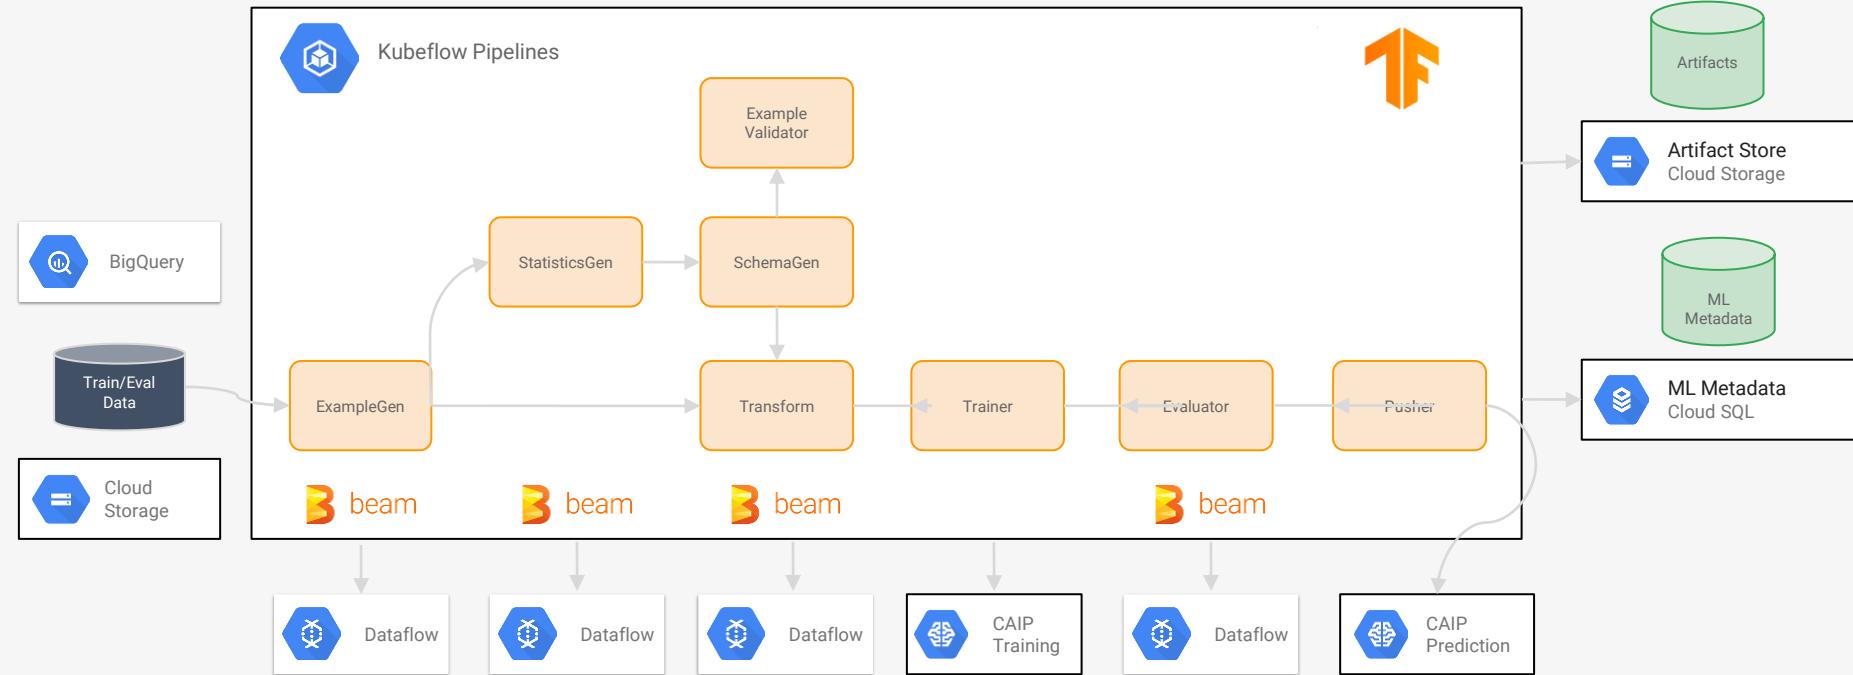
Executors do the work





TFX in Cloud AI Platform

Google Cloud Platform





More Information



TFX Website: <https://www.tensorflow.org/tfx>

TFX Repo: <https://github.com/tensorflow/tfx>

TFX Community: <https://goo.gle/tfx-group>

TFX YouTube: <https://goo.gle/tfx-youtube>



Tutorials

TFX on AI Platform Pipelines: <https://www.tensorflow.org/tfx/tutorials/tfx/cloud-ai-platform-pipelines>

MiniKF on VM Instance: <https://codelabs.developers.google.com/codelabs/cloud-kubeflow-minikf-kale/>

KubeFlow on GKE: <https://codelabs.developers.google.com/codelabs/kubeflow-introduction>

Github Issue Summarization: <https://codelabs.developers.google.com/codelabs/cloud-kubeflow-pipelines-gis/>

Pythian

L♥VE YOUR DATA

Perguntas?

Obrigado!